
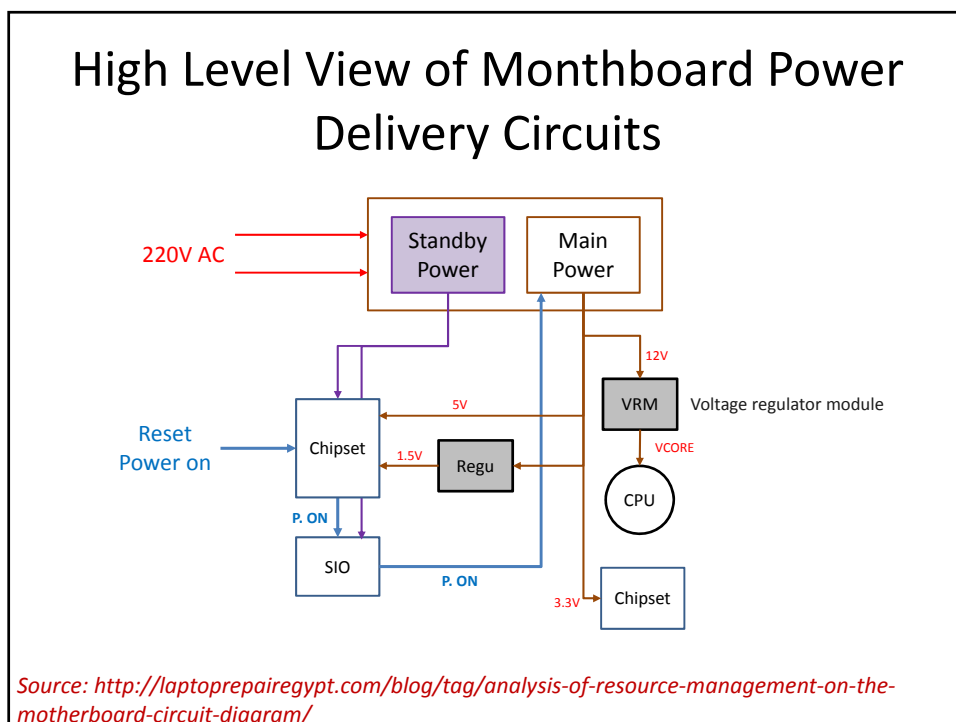
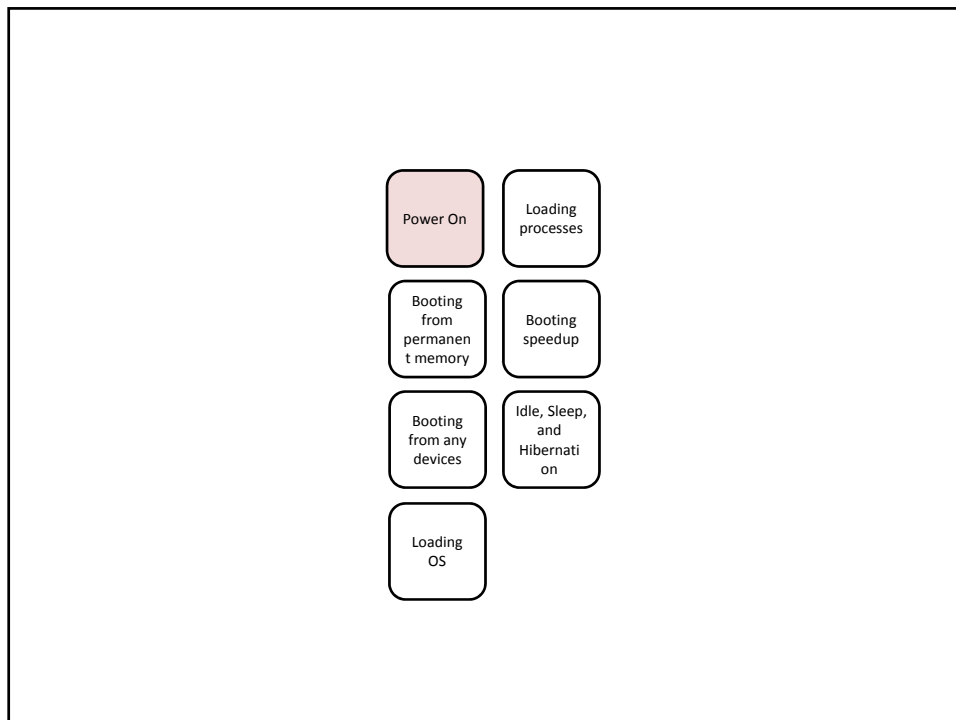


# Operating System Design and Implementation

*Booting process*


Shiao-Li Tsao





## References

- <http://www.kitguru.net/components/power-supplies/ironlaw/how-computer-power-supplies-work-kitguru-guide/>
- <http://www.hardwaresecrets.com/everything-you-need-to-know-about-the-motherboard-voltage-regulator-circuit/>



## Why four loaders are necessary

- First loader (bootloader in permanent memory such as ROM/NOR flash/randomly access and byte addressable)
  - Processor can run but knows nothing about the whole system
  - Prepared by by SoC vendors or motherboard vendors to make sure there is no problem with the whole system
  - Check/initialize the hardware
  - May provide basic services to other programs
  - May provide shells to end-users for basic operations
  - Load/jump to the second loader (based on predefined procedures/configurations)

## Why four loaders are necessary (Cont.)

- Second loader (bootloader in any devices such as DISK/CD-ROM/NAND flash/remote server/...)
  - As soon as first bootloader can access the second loader (on any device), load it into memory, and can jump to the starting instruction
  - Prepare by OS vendors or other third party vendors
    - Stored in the correct location/in correct format)
    - Must know how to loader OSs
  - May provide shell to end-users to select OSs (multi-booting) or set configurations
  - Loading OS into memory
  - Optional for embedded processors

## Why four loaders are necessary (Cont.)


- Third loader (OS loader that loads OS)
  - A program executes without OS services
  - Initial and prepare OS services
- Fourth loader (processes replies on OS to fork/exec other processes)
  - OS is now ready and can provide services
  - Any process calls OS services to fork/execute other processes

## 1<sup>st</sup> boot-loader functions

- Initialize the hardware setting
  - Power on self test (POST) in x86/BIOS
- Basic monitor and debugger
- Pass the control to the 2<sup>nd</sup> bootloader

## First Bootlader - embedded processor

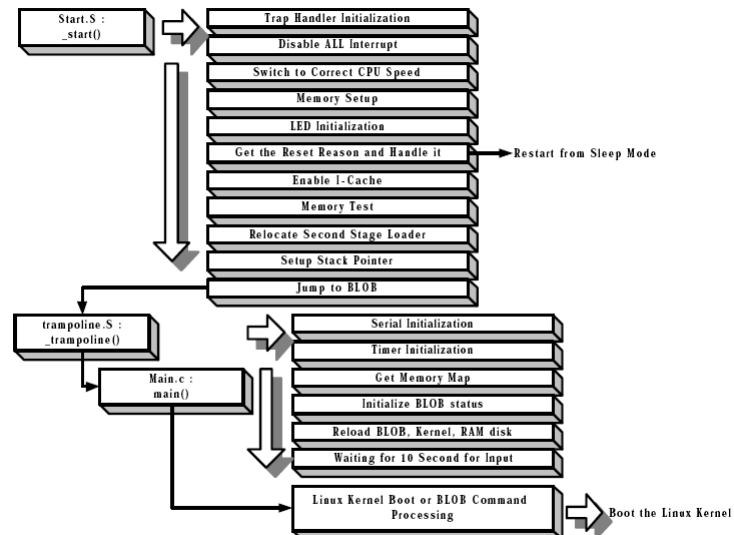
- SoC: e.g. STM32F4
- Reference:
  - <http://www.st.com/web/en/resource/technical/document/datasheet/DM00037051.pdf>
  - p.18



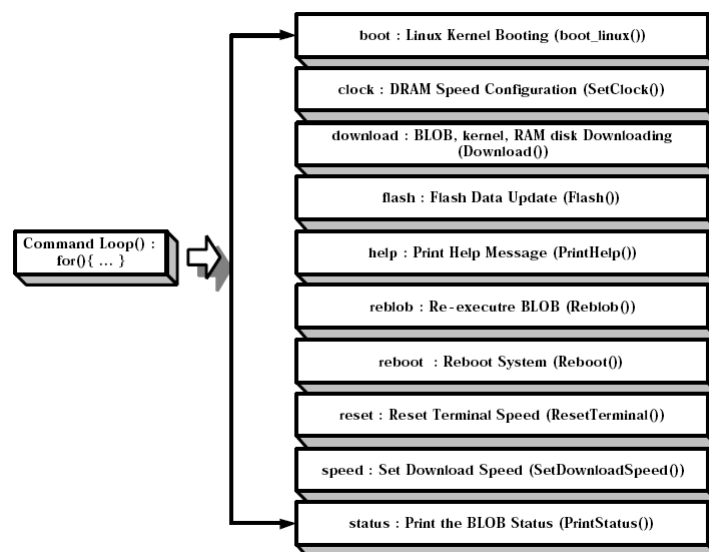
## First Bootlader - embedded processor

- SoC: e.g. STM32F4
  - Memory space
- Reference:
  - <http://www.st.com/web/en/resource/technical/document/datasheet/DM00037051.pdf>
  - p.70

## Boot-loader Example



## Boot-loader Example



# ARM Examples

- Vector table

```
.text
/* Jump vector table as in table 3.1 in [1] */
.globl _start
_start:    b      reset
           b      undefined_instruction
           b      software_interrupt
           b      prefetch_abort
           b      data_abort
           b      not_used
           b      irq
           b      fiq
```

# Reset\_Handler

```
/* the actual reset code */
reset:
    /* First, mask ALL interrupts */
    ldr    r0, IC_BASE
    mov    r1, #0x00
    str    r1, [r0, #ICMR]

    /* switch CPU to correct speed */
    ldr    r0, PWR_BASE
    LDR    r1, cpuspeed
    str    r1, [r0, #PPCR]

    /* setup memory */
    bl     memsetup

    /* init LED */
    bl     ledinit
```



## Reset\_Handler (Cont.)

```
/* check if this is a wake-up from sleep */
ldr    r0, RST_BASE
ldr    r1, [r0, #RCSR]    Reset status register
and    r1, r1, #0x0f
teq    r1, #0x08
bne    normal_boot    /* no, continue booting */

/* yes, a wake-up. clear RCSR by writing a 1 (see
9.6.2.1 from [1]) */
mov    r1, #0x08
str    r1, [r0, #RCSR]    ;

/* get the value from the PSPR and jump to it */
ldr    r0, PWR_BASE
ldr    r1, [r0, #PSPR]    Power manager scratch pad register
mov    pc, r1
```

## Reset\_Handler (Cont.)

```
normal_boot:
/* enable I-cache */
mrc    p15, 0, r1, c1, c0, 0    @ read control reg
orr    r1, r1, #0x1000    @ set Icache
mcr    p15, 0, r1, c1, c0, 0    @ write it back

/* check the first 1MB in increments of 4k */
mov    r7, #0x1000
mov    r6, r7, lsl #8    /* 4k << 2^8 = 1MB */
ldr    r5, MEM_START

mem_test_loop:
mov    r0, r5
bl     testram
teq    r0, #1
beq    badram

add    r5, r5, r7
subs   r6, r6, r7
bne    mem_test_loop
```

```

/* the first megabyte is OK, so let's clear it */
mov    r0, #((1024 * 1024) / (8 * 4)) /* 1MB in
steps of 32 bytes */
ldr     r1, MEM_START
...
clear_loop:
    stmia r1!, {r2-r9}
    subs  r0, r0, #(8 * 4)
    bne   clear_loop

/* relocate the second stage loader */
add     r2, r0, #(128 * 1024) /* blob is 128kB
*/
add     r0, r0, #0x400 /* skip first 1024
bytes */
ldr     r1, MEM_START
add     r1, r1, #0x400 /* skip over here
as well */
..
copy_loop:
    ldmia r0!, {r3-r10}
    stmia r1!, {r3-r10}
    cmp   r0, r2
    ble   copy_loop

```

## Reset\_Handler (Cont.)


```

/* set up the stack pointer */
ldr     r0, MEM_START
add     r1, r0, #(1024 * 1024)
sub     sp, r1, #0x04

/* blob is copied to ram, so jump to it */
add     r0, r0, #0x400
mov     pc, r0

```

## Bootloader memory map



## Boot-loader C program

```
int main(void)
{
~
    led_on();
~
    SerialInit(baud9k6);
    TimerInit();
~
    SerialOutputString(PACKAGE " version " VERSION "\n"
        "Copyright (C) 1999 2000 2001 ");
~
    get_memory_map();
~
    SerialOutputString("Running from ");
    if(RunningFromInternal())
        SerialOutputString("internal");
    else
        SerialOutputString("external");
~
    ...
}
```

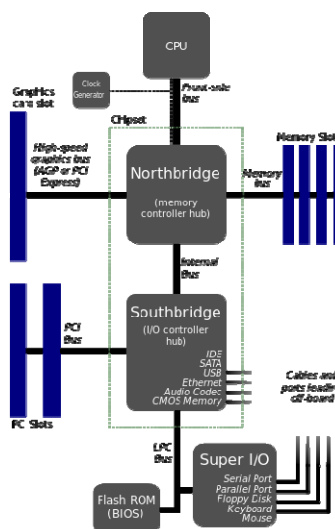
## Boot-loader C program

```

/* wait 10 seconds before starting autoboot */
SerialOutputString("Autoboot in progress, press any key...");
for(i = 0; i < 10; i++) {
    SerialOutputByte('.');
    retval = SerialInputBlock(commandline, 1, 1);
~
    if(retval == 0) {
~
        boot_linux(commandline);    }
~
    for(;;) {
~
        if(numRead > 0) {
            if(MyStrNCmp(commandline, "boot", 4) == 0) {
                boot_linux(commandline + 4);
            } else if(MyStrNCmp(commandline, "clock", 5) == 0) {
                SetClock(commandline + 5);
            } else if(MyStrNCmp(commandline, "download ", 9) == 0) {
~
            return 0;
        } /* main */
    }

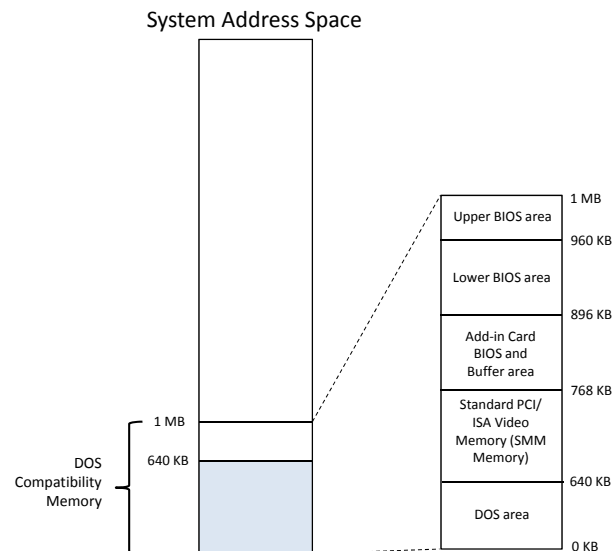
```

## First Bootloader – x86

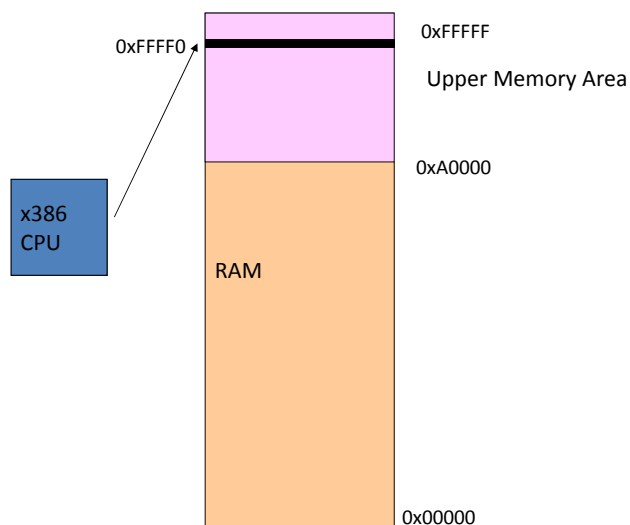


Source: Wikipedia ([https://en.wikipedia.org/wiki/Northbridge\\_\(computing\)\)](https://en.wikipedia.org/wiki/Northbridge_(computing)))

## First Bootloader – x86



## PC Booting (Cont)




# UMA

Address	First 16K (x0000h- x3FFFh)	Second 16K (x4000h- x7FFFh)	Third 16K (x8000h- xBFFFh)	Fourth 16K (xC000h- xFFFFh)
<b>A0000-AFFFFh</b>	VGA Graphics Mode Video			
<b>B0000- BFFFFh</b>	VGA Monochrome Text Mode Video RAM		VGA Color Text Mode Video RAM	
<b>C0000- CFFFFh</b>	VGA Video BIOS ROM		IDE Hard Disk BIOS ROM	Optional Adapter ROM BIOS or RAM UMBs
<b>D0000- DFFFFh</b>	Optional Adapter ROM BIOS or RAM UMBs			
<b>F0000- FFFFFh</b>	System BIOS ROM			

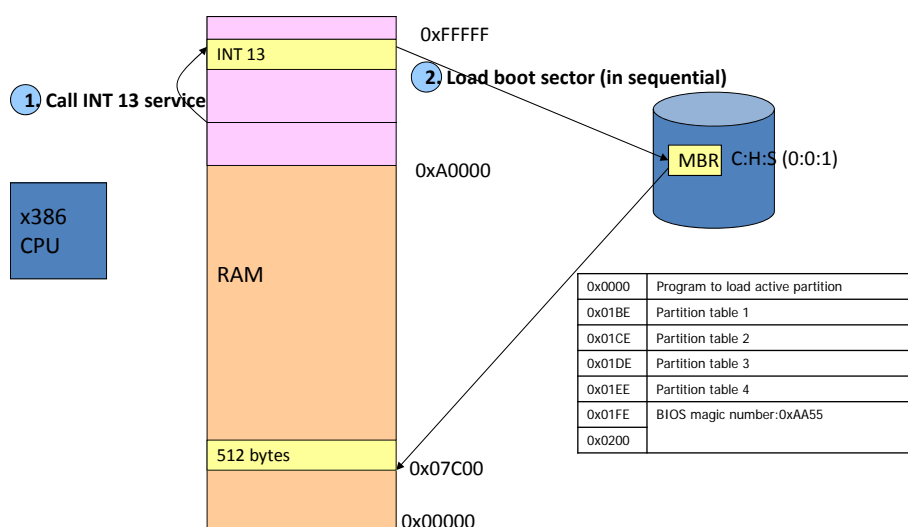
## PC Booting (Cont)

1. Power supply sends POWER GOOD to CPU
2. CPU resets
3. Run FFFF:0000 @ BIOS ROM
4. Jump to a real BIOS start address
5. POST
6. Beep if there is an error
7. Read CMOS data/settings
8. Run 2<sup>nd</sup>-stage boot

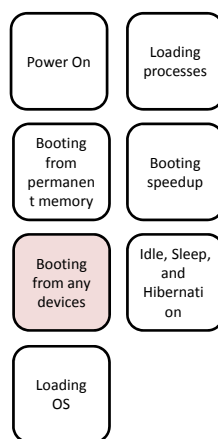



INT	Address	Type	Description
00h	0000:0000h	Processor	Divide by zero
01h	0000:0004h	Processor	Single step
02h	0000:0008h	Processor	Non maskable interrupt (NMI)
03h	0000:000Ch	Processor	Breakpoint
04h	0000:0010h	Processor	Arithmetic overflow
05h	0000:0014h	Software	Print screen
06h	0000:0018h	Processor	Invalid op code
07h	0000:001Ch	Processor	Coprocessor not available
08h	0000:0020h	Hardware	System timer service routine
09h	0000:0024h	Hardware	Keyboard device service routine
0Ah	0000:0028h	Hardware	Cascade from 2nd programmable interrupt controller
0Bh	0000:002Ch	Hardware	Serial port service - COM post 2
0Ch	0000:0030h	Hardware	Serial port service - COM port 1
0Dh	0000:0034h	Hardware	Parallel printer service - LPT 2
0Eh	0000:0038h	Hardware	Floppy disk service
0Fh	0000:003Ch	Hardware	Parallel printer service - LPT 1
10h	0000:0040h	Software	Video service routine
11h	0000:0044h	Software	Equipment list service routine
12h	0000:0048H	Software	Memory size service routine
13h	0000:004Ch	Software	Hard disk drive service
14h	0000:0050h	Software	Serial communications service routines
15h	0000:0054h	Software	System services support routines
16h	0000:0058h	Software	Keyboard support service routines
17h	0000:005Ch	Software	Parallel printer support services

## PC Booting (Cont)

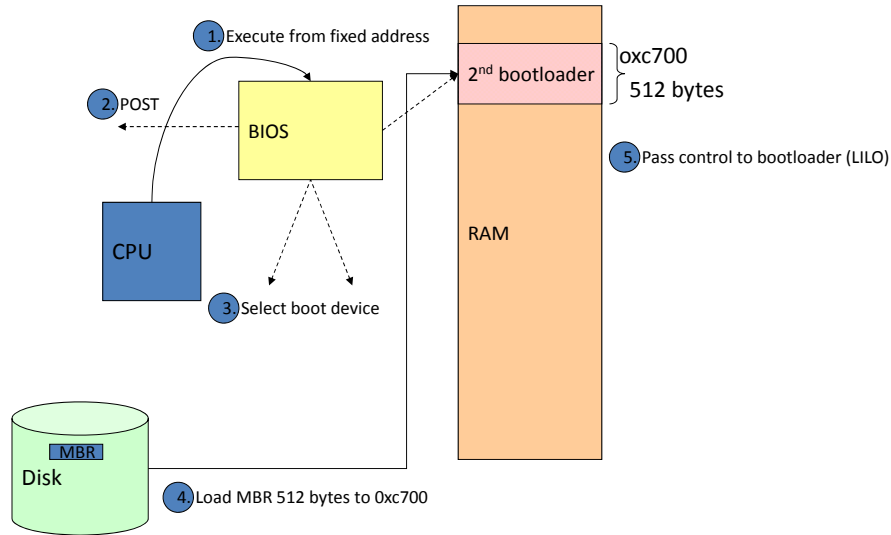


# MBR

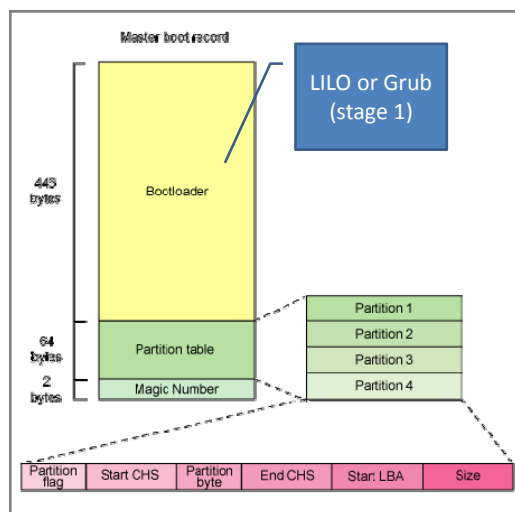




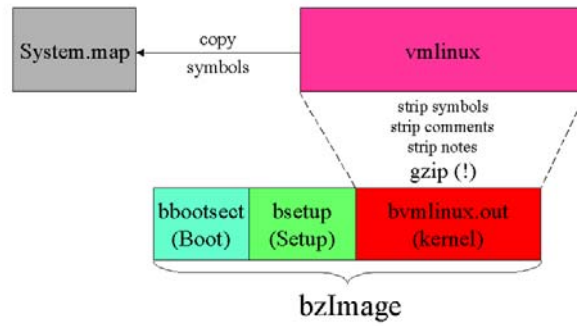
# Linux Boot Example




## MBR (Master Boot Record)



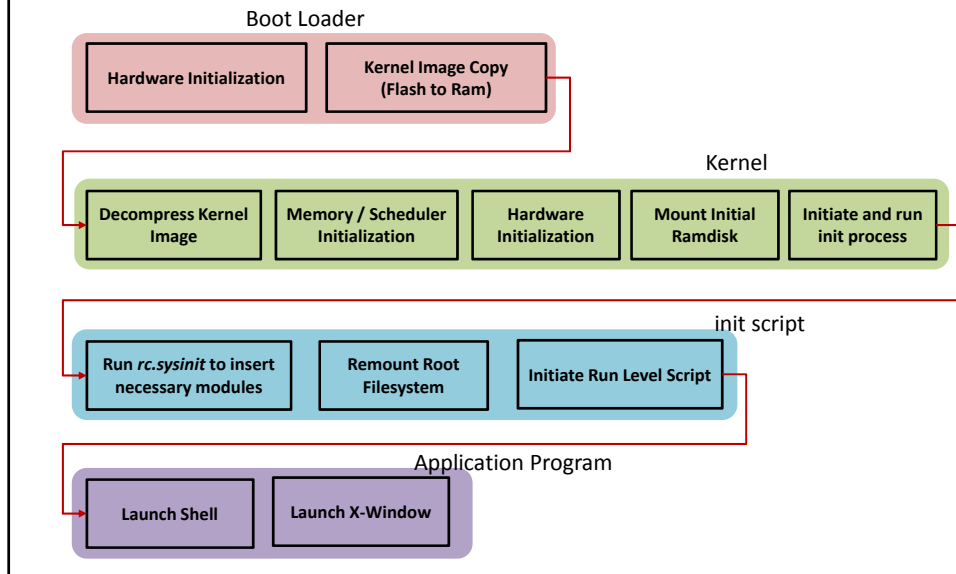
## Anatomy of bzImage



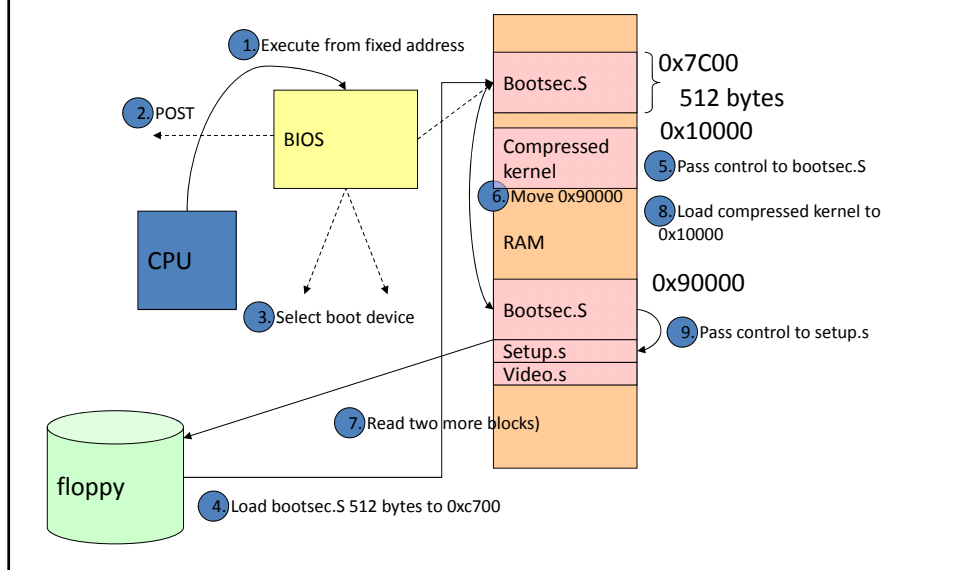
Source: <https://commons.wikimedia.org/wiki/File:Anatomy-of-bzimage.png>



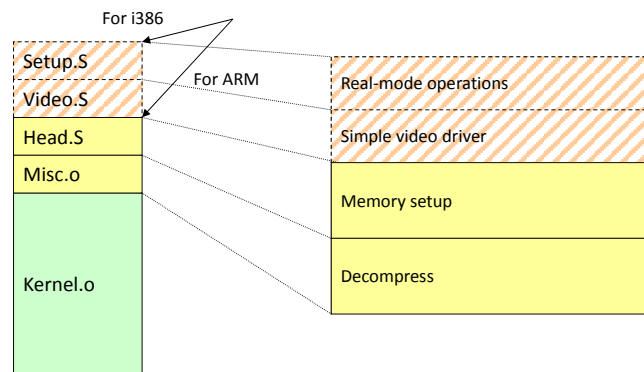
# Loading Kernel



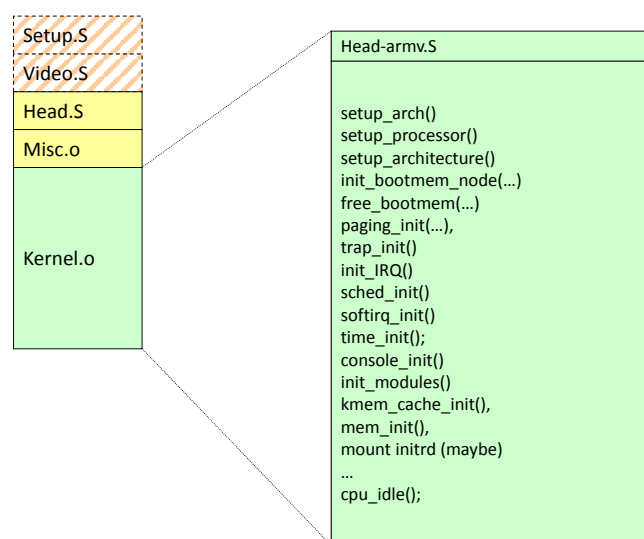
# Linux Boot Example



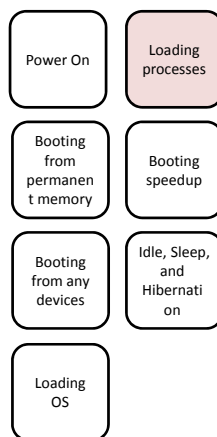
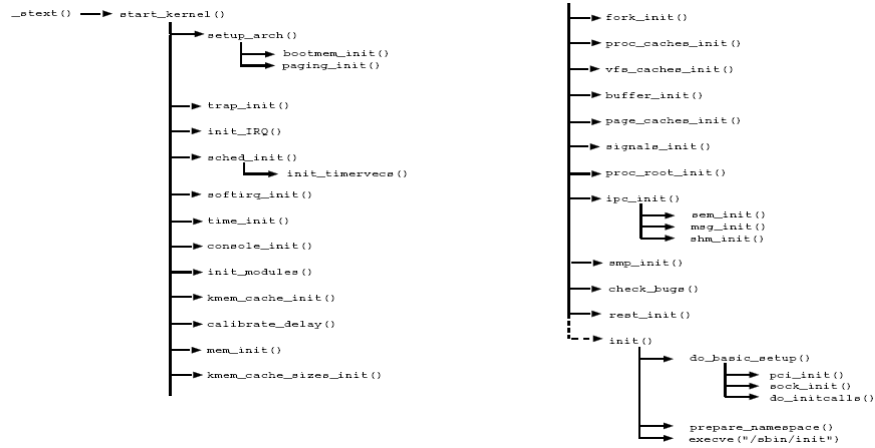
# Kernel Image Structure



# Kernel Image Structure



# Linux booting

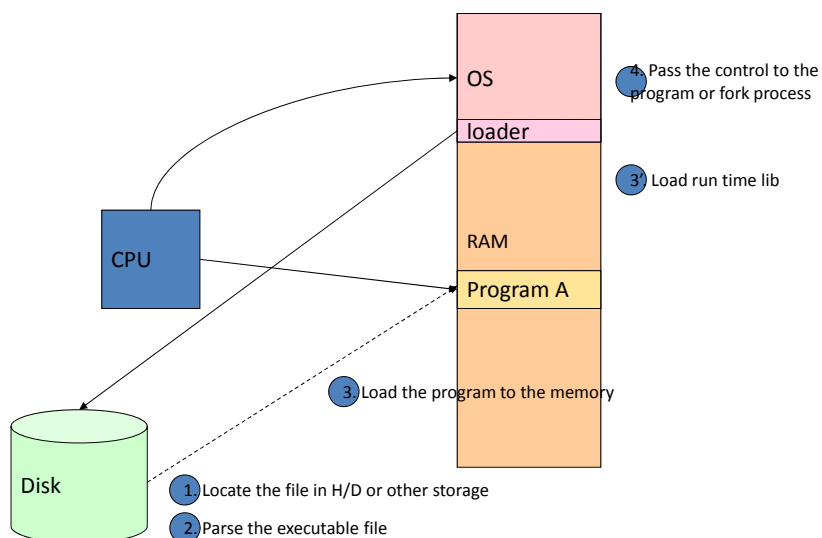


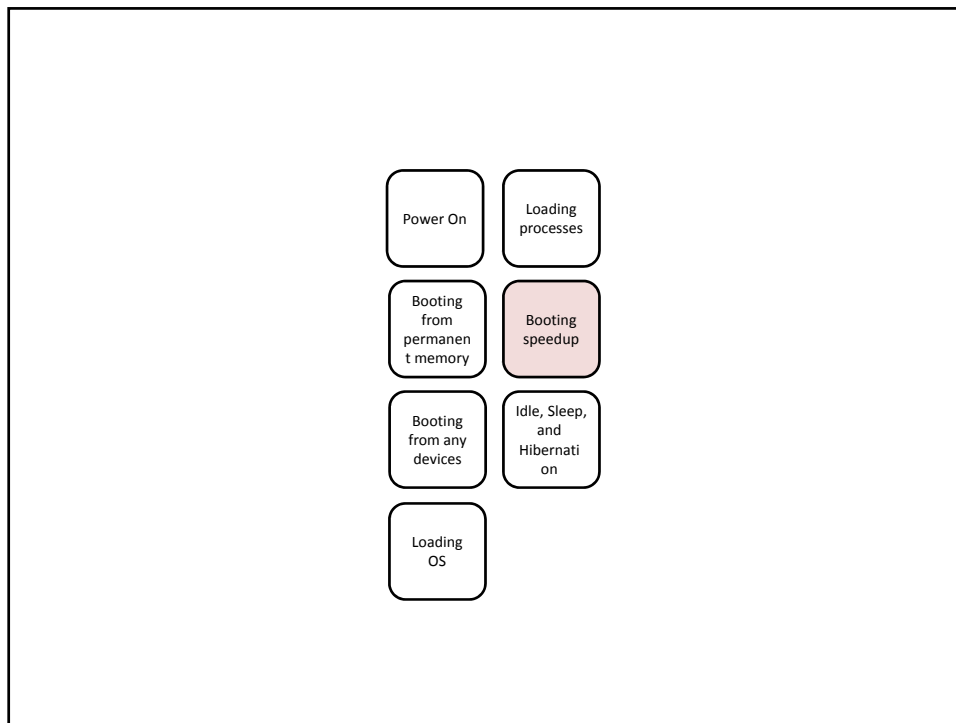
# Mount root disk

Source: <http://linux-development-for-fresher.blogspot.tw/2012/07/linux-boot-process-in-nutshell.html>

initrd (initramfs): boot loader initialized RAM disk

## Loader





## Bootloader Speedups

- Remove waiting time
- Removing unnecessary initialization routines
- Uncompressed kernel
- DMA Copy Of Kernel On Startup
- Fast Kernel Decompression
- Kernel XIP

## Uncompressed kernel

- Fast boot, but image size has been larger
  - 2MB – 2.5MB non-compressed (ARM)
  - 1MB – 1.5MB compressed (ARM)
- It should be different results, performance of CPUs, speed of flash memory
  - Profiling is required
- Make Image vs. make zImage

## Fast Kernel Decompression

- Use other compression/decompression algorithm
  - Slow compression, good compression ratio, fast decompression
- GZIP vs. Sony UCL
- Small kernel size, fast kernel loading time



## Fast Kernel Decompression

Image file:	initrd-2.6.5-1.358		Power PC
method	UCL	GZIP	improved %
parameter	-b4194304	-8	.
source file size	819200	819200	.
compressed size	187853	189447	.
compression rate	77.1%	76.9%	0.3%
compression time: user (sec)	5.13	2.03	-152.5%
sys (sec)	0.09	0.06	-36.5%
total (sec)	5.22	2.09	-149.0%
decompression time: user (sec)	0.12	0.3	59.7%
sys (sec)	0.1	0.08	-16.9%
total (sec)	0.22	0.39	43.0%

## Fast Kernel Decompression

Image file:	vmlinux-2.4.20 for ibm-440gp		PowerPC
method	UCL	GZIP	improved %
parameter	-b4194304	-8	.
source file size	1810351	1810351	.
compressed size	790250	776807	.
compression rate	56.3%	57.1%	-1.3%
compression time: user (sec)	17.29	6.07	-185.0%
sys (sec)	0.04	0.02	-92.4%
total (sec)	17.33	6.09	-184.6%
decompression time: user (sec)	0.12	0.16	26.1%
sys (sec)	0.03	0.04	35.8%
total (sec)	0.15	0.2	28.2%

## Kernel XIP

- Direct addressing on NOR flash memory
- Cannot compress kernel
- PowerPC 405LP/266 MHZ

Boot Stage	Non-XIP Time	XIP Time
Copy kernel to RAM	85 ms	12 ms *
Decompress kernel	453 ms	0 ms
Kernel time to initialize (time to first user space program)	819 ms	882 ms
Total kernel boot time	1357 ms	894 ms
<b>Reduction:</b>	*	<b>463 ms</b>

## Kernel XIP (Cont.)

- TI OMAP 5912/196 MHZ

Boot Stage	Non-XIP Time Kernel compressed	Non-XIP Time Kernel not compressed	XIP Time
Copy kernel to RAM	56 ms	120 ms	0 ms
Decompress kernel	545 ms	0 ms	0 ms
Kernel time to initialize (time to first user space program)	88 ms	208 ms	110 ms
Total kernel boot time	689 ms	208 ms	110 ms
<b>Reduction:</b>	*	<b>481 ms</b>	<b>579 ms</b>

## Embedded Linux Speedups

- Removing unnecessary message printout
- Remove unnecessary functions and device drivers
- Modularization of device driver
- Asynchronous function calls
- Avoid performance measurement routine
- RTC no sync
- Using read-only file system
- Using lazy mount technique on R/W file systems
- Deferred Initcalls

## Application Speedups

- Using binary script, not shell script
- Using init process with simplified and optimized
- Parallel RC scripts
- Application XIP
- Using pre-link shared lib
- Optimize of application programs
- Move from glibc to uClibc