# Operating System Design and Implementation

*Process Management – Part II*

Shiao-Li Tsao

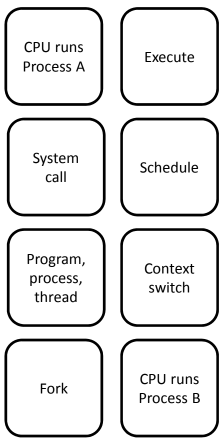

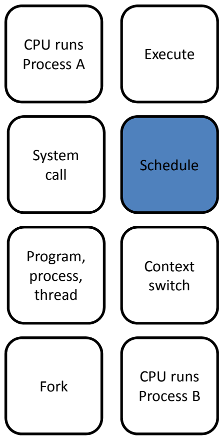


## Process schedule and context switching in Linux

- Scheduling
  - Find the next suitable process to run
- Context switch
  - Store the context of the current process, restore the context of the next process

# Scheduler in Details

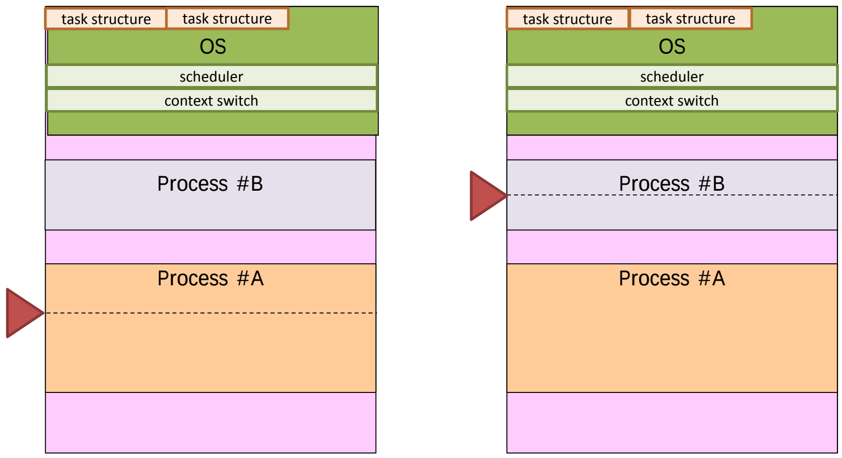


# Process schedule and context switching in Linux

- When is the scheduler be invoked
  - Direct invocation vs. Lazy invocation

  - When returning to user-space from a system call
  - When returning to user-space from an interrupt handler
  - When an interrupt handler exits, before returning to kernel-space
  - If a task in the kernel explicitly calls schedule()
  - If a task in the kernel blocks (which results in a call to schedule())

6

# Interrupt Basics

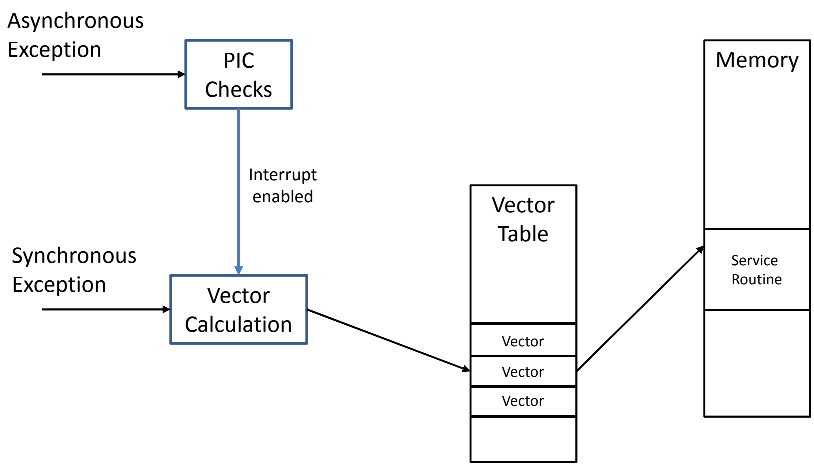


*Source: Qing Li "real-time concepts for embedded systems"*

# X86 Interrupts

| Vector range | Use |
|---|---|
| 0x0 – 0x13 | Non-maskable interrupts/exceptions |
| 0x14 – 0x1F | Intel-reserved |
| 0x20 – 0x7F | External interrupts (IRQs) |
| 0x80 | Programmed exception for system calls |
| 0x81 – 0xEE | External interrupts (IRQs) |

| IRQ | Function |
|---|---|
| 0 | System Timer |
| 1 | Keyboard Controller |
| 2 | 2nd IRQ Controller Cascade |
| 8 | Real-Time Clock |
| 9 | Avail. |
| 10 | Available |
| 11 | Available |
| 12 | Mouse Port / Available |
| 13 | Math Coprocessor |
| 14 | Primary IDE |
| 15 | Secondary IDE |
| 3 | Serial 2 |
| 4 | Serial 1 |
| 5 | Sound card / Parallel 2 |
| 6 | Floppy Disk Controller |
| 7 | Parallel 1 |

*Source: https://en.wikipedia.org/wiki/Interrupt_request_(PC_architecture)*

# Time Interrupt Basics

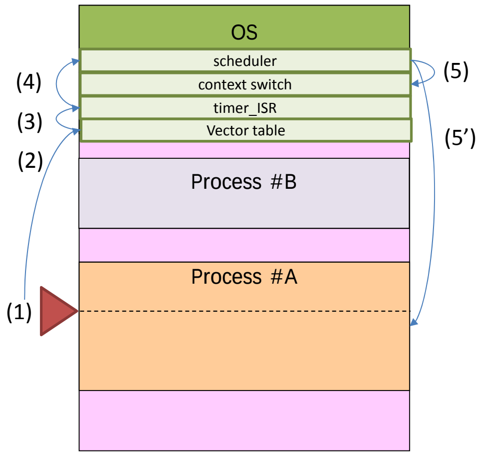


# Process Context

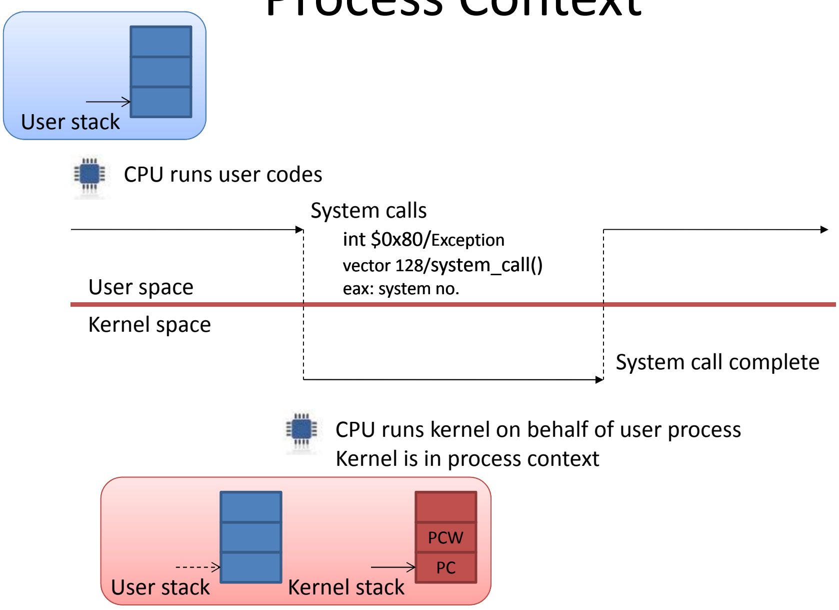

# Interrupt Context

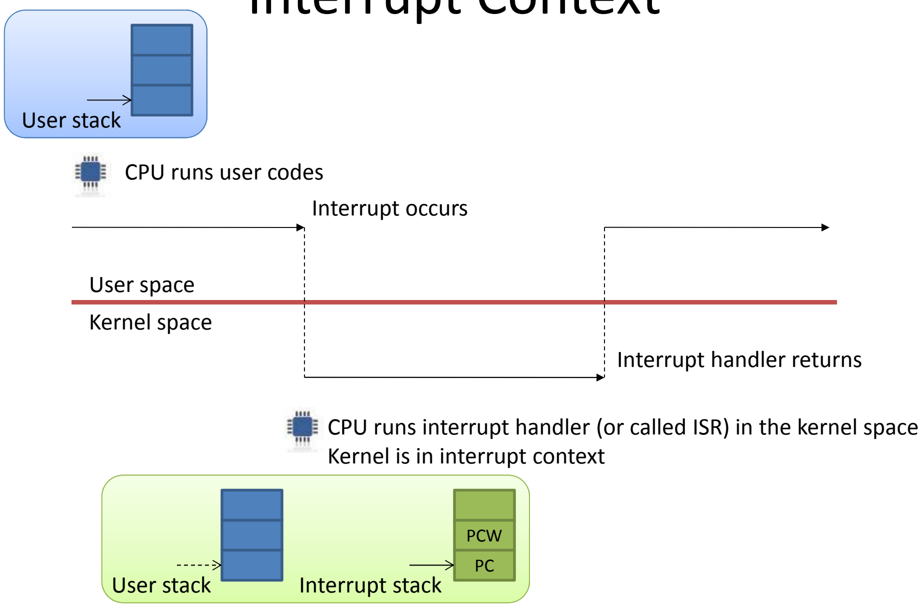


# Interrupt Context

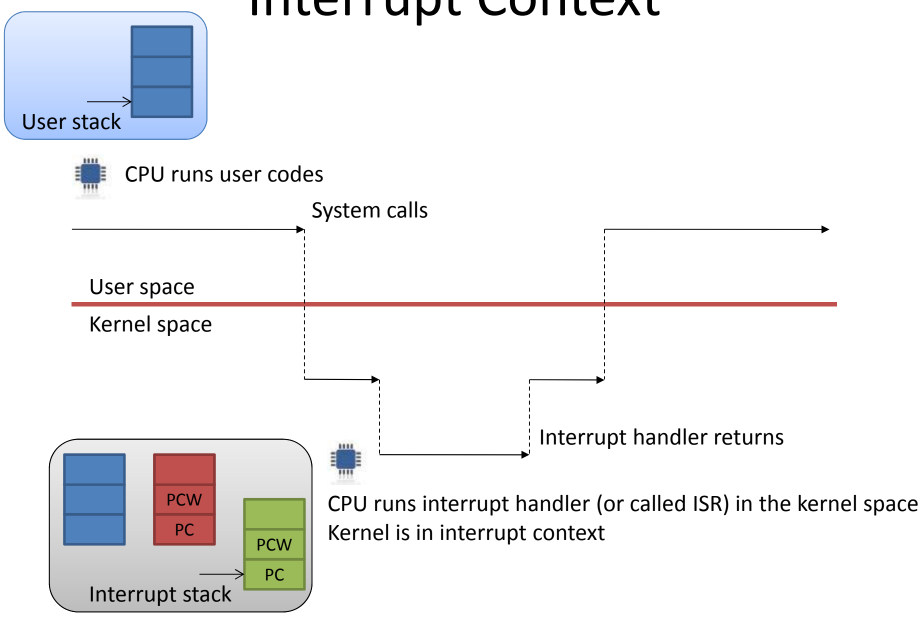

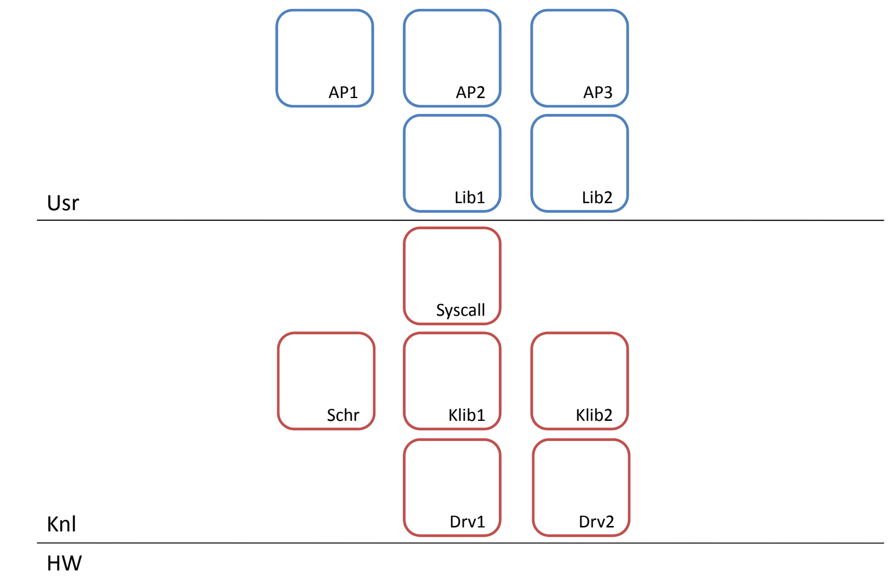


When returning to user-space from a system call

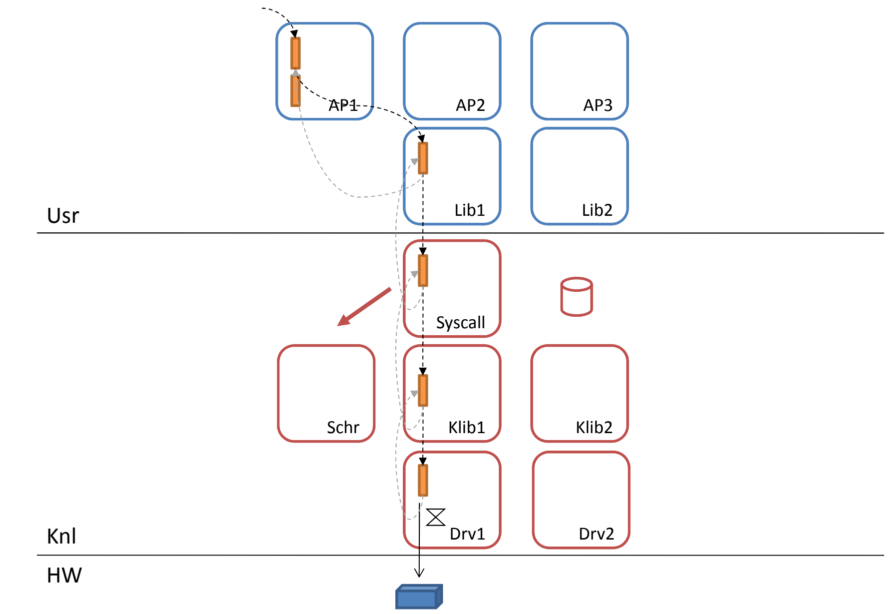

If a task in the kernel blocks (which results in a call to schedule())

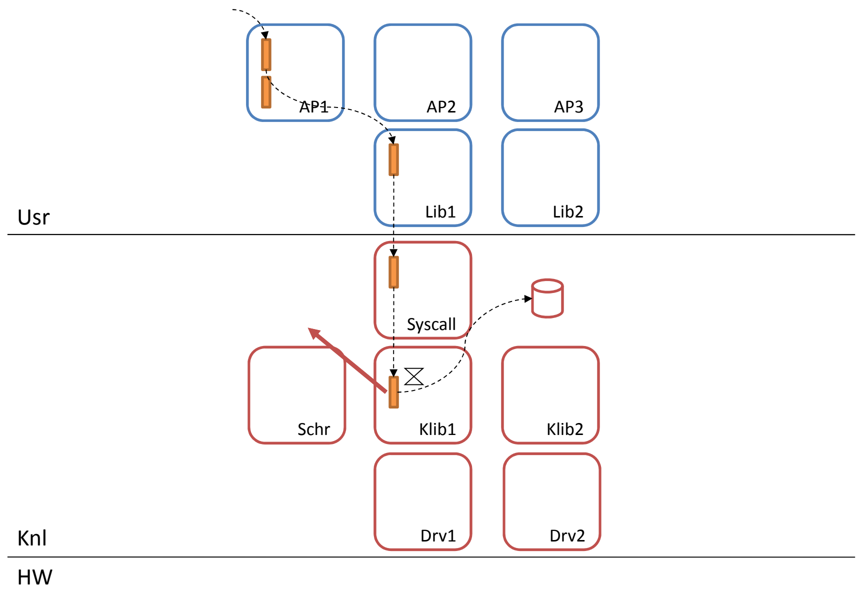


When returning to user-space from an interrupt handler

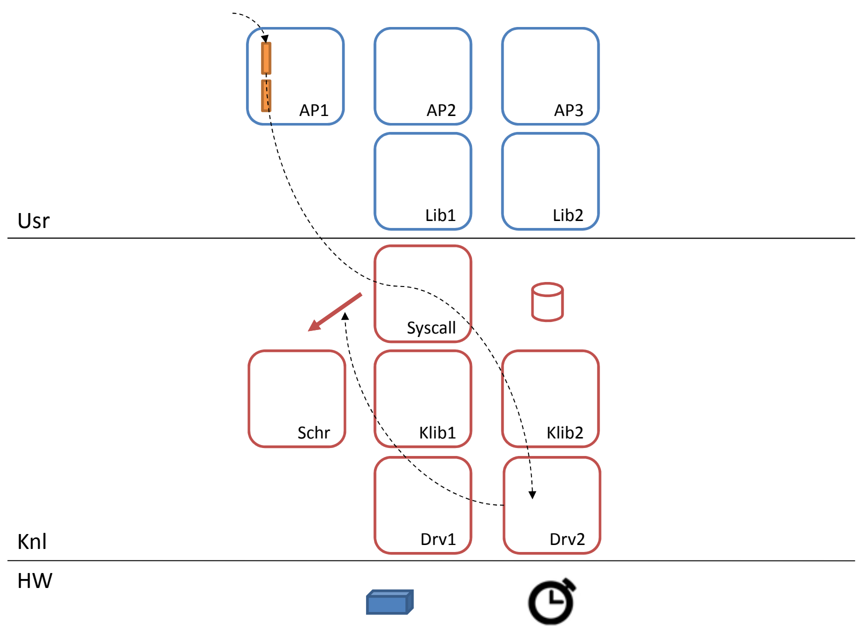

When returning to user-space from an interrupt handler

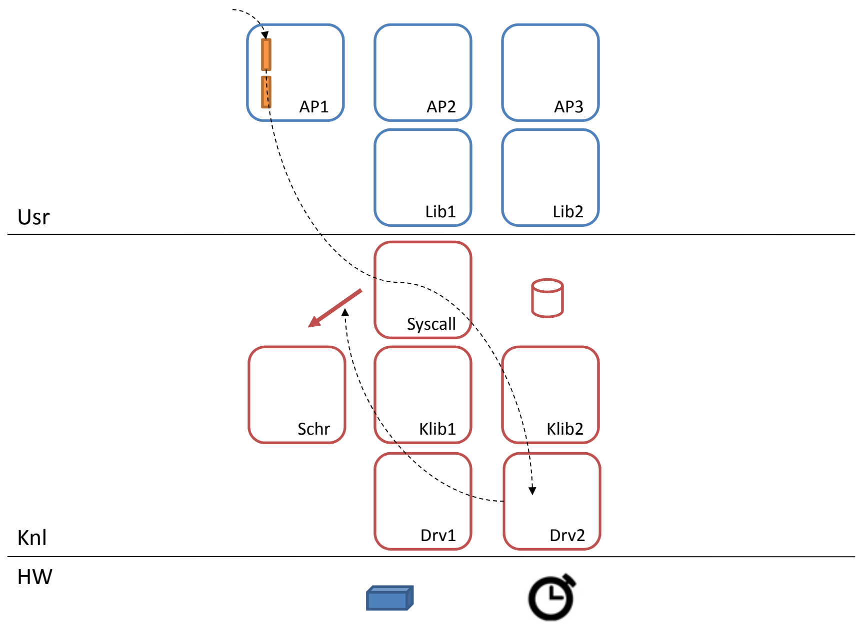


When an interrupt handler exits, before returning to kernel-space

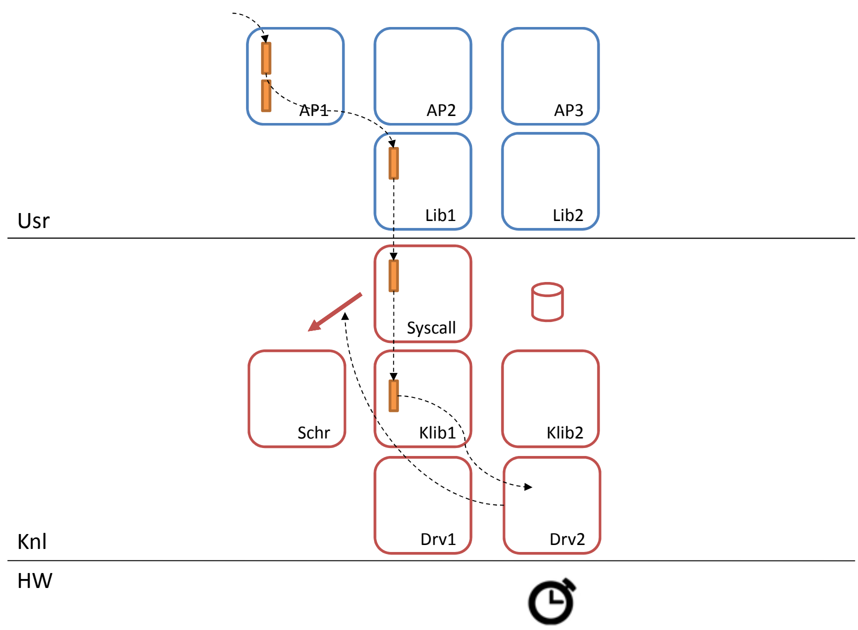

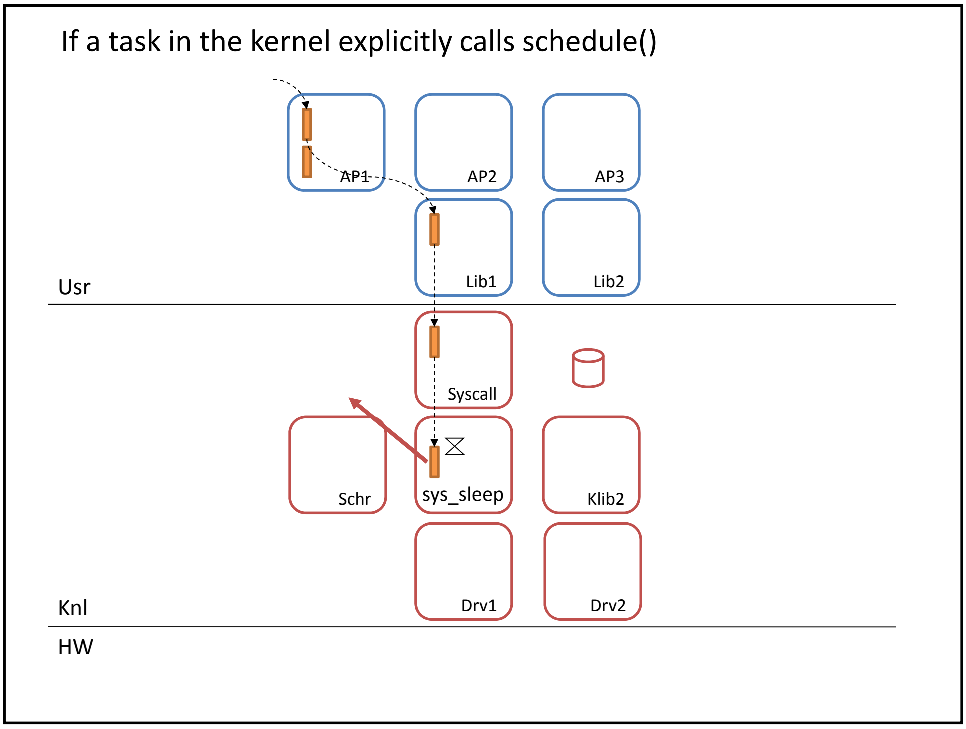
If a task in the kernel explicitly calls schedule()
AP1
AP2
AP3
Lib1
Lib2
Usr
Syscall
Schr
sys_sleep
Klib2
Drv1
Drv2
Knl
HW

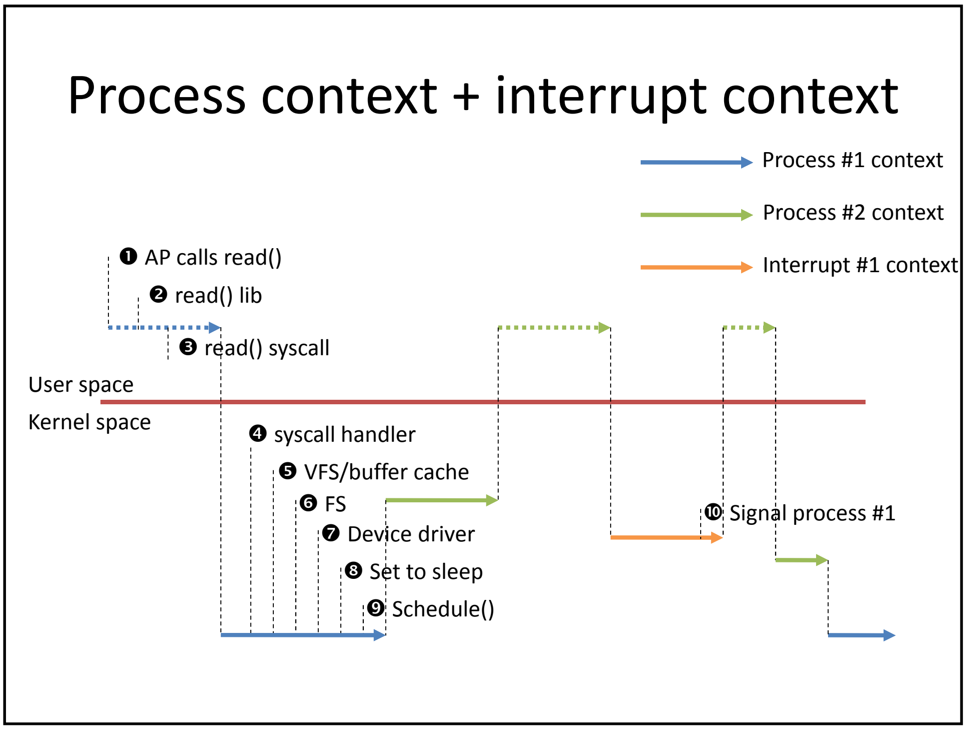
Process context + interrupt context
Process #1 context
Process #2 context
Interrupt #1 context
❶ AP calls read()
❷ read() lib
❸ read() syscall
User space
Kernel space
❹ syscall handler
❺ VFS/buffer cache
❻ FS
❼ Device driver
❽ Set to sleep
❾ Schedule()
❿ Signal process #1

# User preemption

- User preemption occurs when the kernel is in a safe state and about to return to user-space

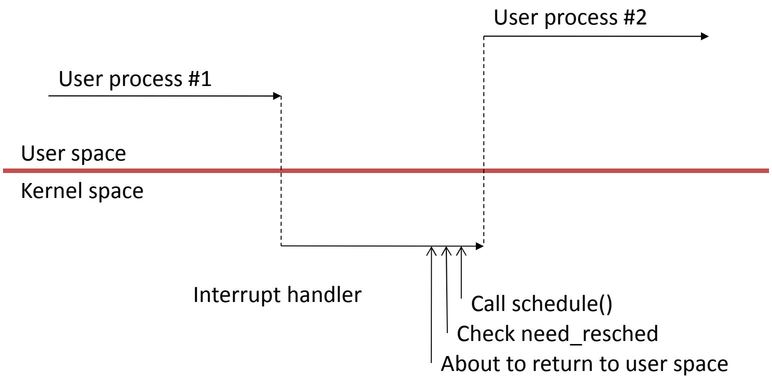



# User preemption

- User preemption occurs when the kernel is in a safe state and about to return to user-space


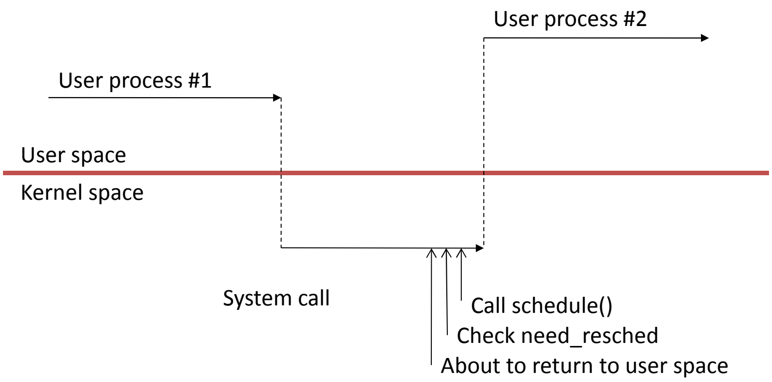

# Kernel preemption

- Linux kernel is possible to preempt a task at any point, so long as the kernel does not hold a lock

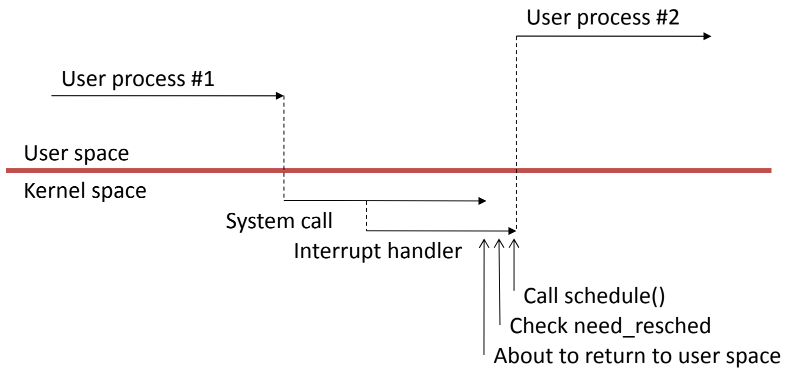


# Preemptive Kernel

- Non-preemptive kernel supports user preemption
- Preemptive kernel supports kernel/user preemption

- Kernel can be interrupted ≠ kernel is preemptive
  - Non-preemptive kernel, interrupt returns to interrupted process
  - Preemptive kernel, interrupt returns to any schedulable process

# Preemptive Kernel

- 2.4 is a non-preemptive kernel
- 2.6 is a preemptive kernel
- 2.6 could disable CONFIG_PREEMPT

# Preemptive Kernel

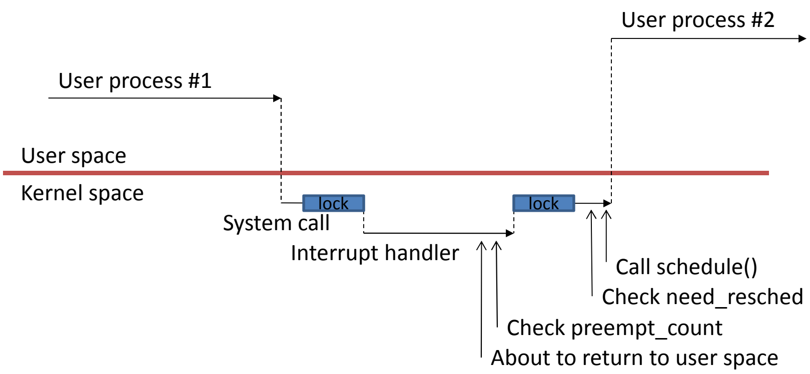

# Preemptive Kernel

- How difficult to implement a preemptive kernel?

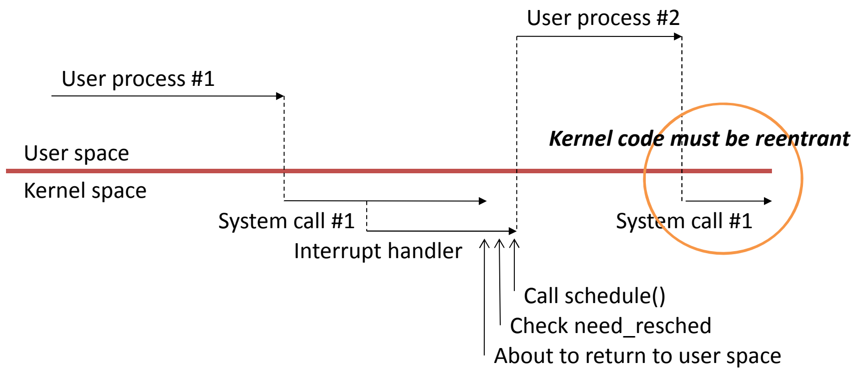


# Single Core vs. Multi-core

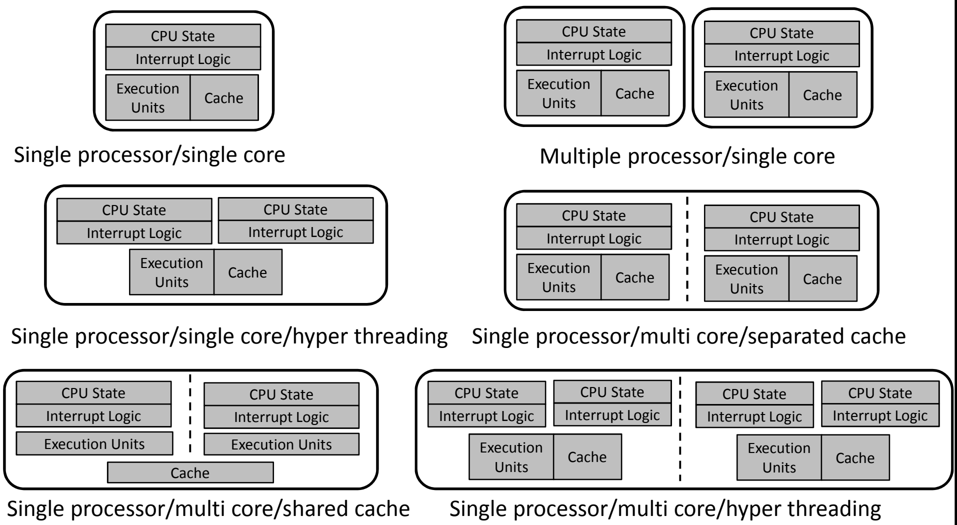

## Single Core vs. Multi-core

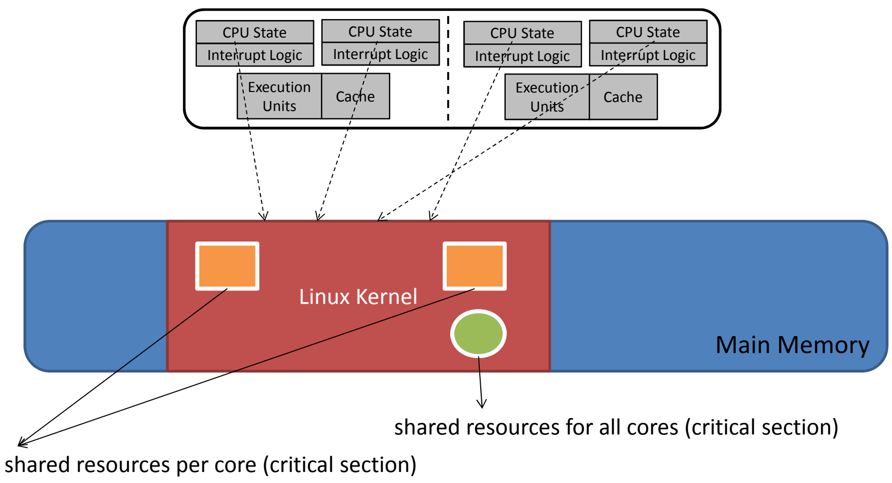


## Schedule Algorithms

- Think about yourself (your homework schedule)
  - Given that you have a lot of homework to do, each with a deadline
  - Profs. continue assigning new homework

  - What is the next homework to do? (next task to schedule)
  - Why should we stop a homework? (time to schedule)
  - How long can we concentrate on a homework? (scheduling period)
  - How long do we spend to determine the next homework? (scheduling algorithm overhead)
  - How much effort do we spend to switch homework? (context switch overhead)
  - What is the importance of a homework? (priority of a job)
  - How long does a homework need? (job length)

# How Linux Scheduler Works

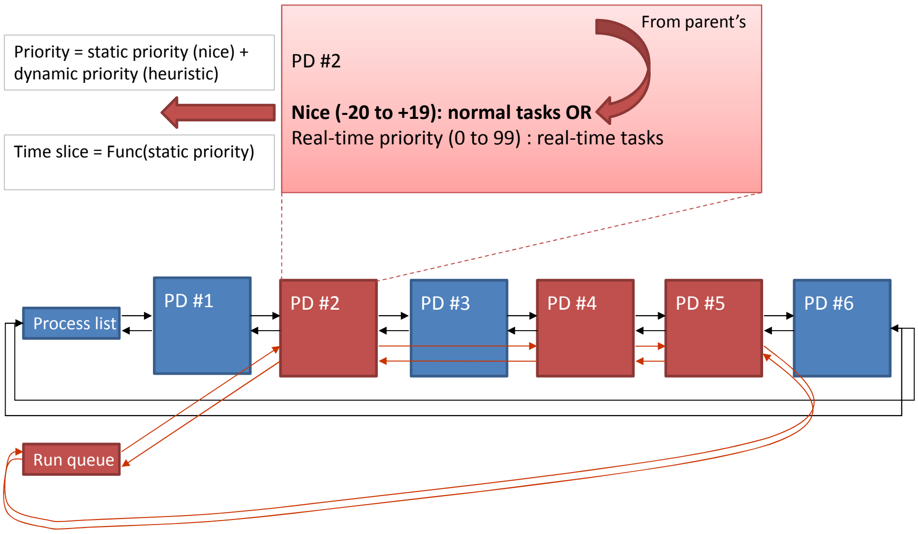


# Timeslice

- Timeslice function

| Type of Task | Nice Value | Timeslice Duration |
|---|---|---|
| Initially created | parent's | half of parent's |
| Minimum Priority | +19 | 5ms (MIN_TIMESLICE) |
| Default Priority | 0 | 100ms (DEF_TIMESLICE) |
| Maximum Priority | -20 | 800ms (MAX_TIMESLICE) |

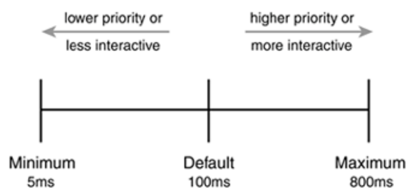


$$\begin{array}{c}\textit{base time quantum}\\ \text{(in milliseconds)}\end{array} = \begin{cases}(140 - \textit{static priority}) \times 20 & \text{if } \textit{static priority} < 120\\ (140 - \textit{static priority}) \times 5 & \text{if } \textit{static priority} \geq 120\end{cases} \quad (1)$$

## Process schedule and context switching in Linux

- Priority-based scheduler
- Dynamic priority-based scheduling
  - Dynamic priority
    - Normal process
      - nice value: -20 to +19 (larger nice values imply you are being nice to others)
  - Static priority
    - Real-time process
      - 0 to 99
  - Total priority: 140

33

## Linux O(1) cheduler

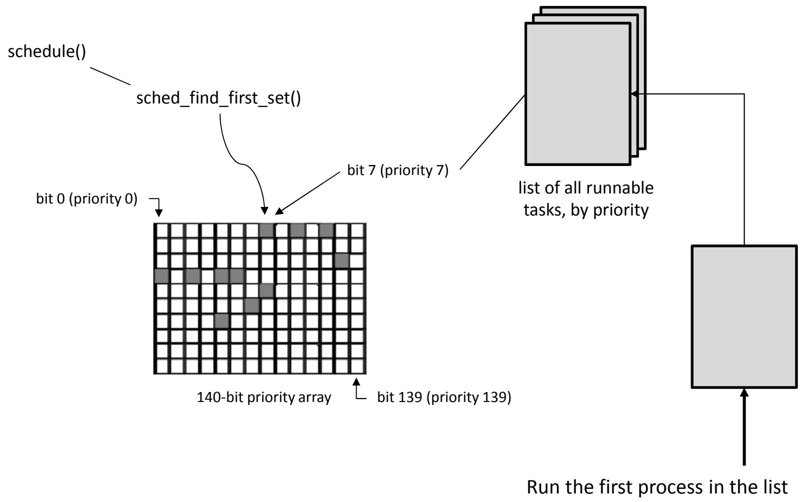

# Linux O(1) cheduler

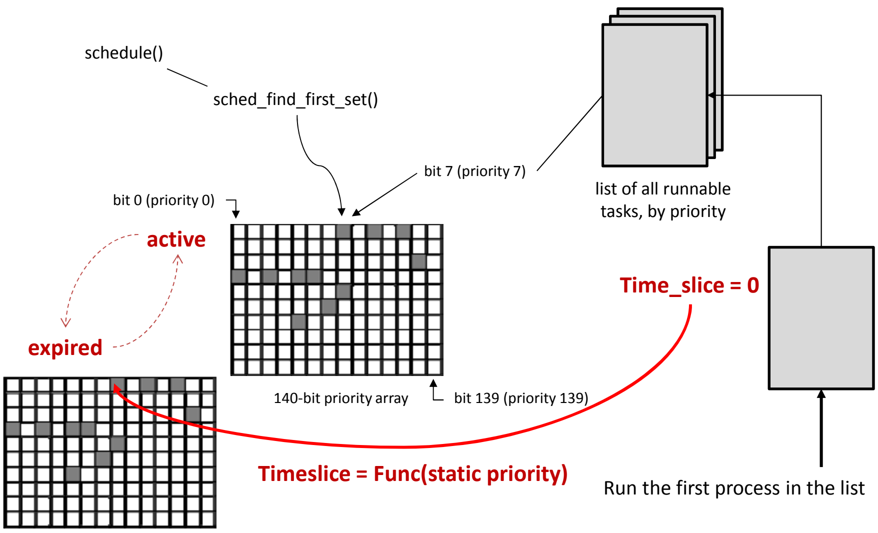


# Calculating Priority

- static_prio = nice
- Prio = nice – bonus + 5

  dynamic priority = max (100, min ( static priority - bonus + 5, 139))

- Heuristic
  – sleep_avg: (0 to MAX_SLEEP_AVG(10ms))
  – sleep_avg+=sleep (becomes runnable)
  – Sleep_avg-=run (every time tick when task runs)

# System calls related to scheduling

| System call | Description |
|---|---|
| nice( ) | Change the static priority of a conventional process |
| getpriority( ) | Get the maximum static priority of a group of conventional processes |
| setpriority( ) | Set the static priority of a group of conventional processes |
| sched_getscheduler( ) | Get the scheduling policy of a process |
| sched_setscheduler( ) | Set the scheduling policy and the real-time priority of a process |
| sched_getparam( ) | Get the real-time priority of a process |
| sched_setparam( ) | Set the real-time priority of a process |
| sched_yield( ) | Relinquish the processor voluntarily without blocking |
| sched_get_ priority_min( ) | Get the minimum real-time priority value for a policy |
| sched_get_ priority_max( ) | Get the maximum real-time priority value for a policy |
| sched_rr_get_interval( ) | Get the time quantum value for the Round Robin policy |
| sched_setaffinity( ) | Set the CPU affinity mask of a process |
| sched_getaffinity( ) | Get the CPU affinity mask of a process |

# How Linux Scheduler Works

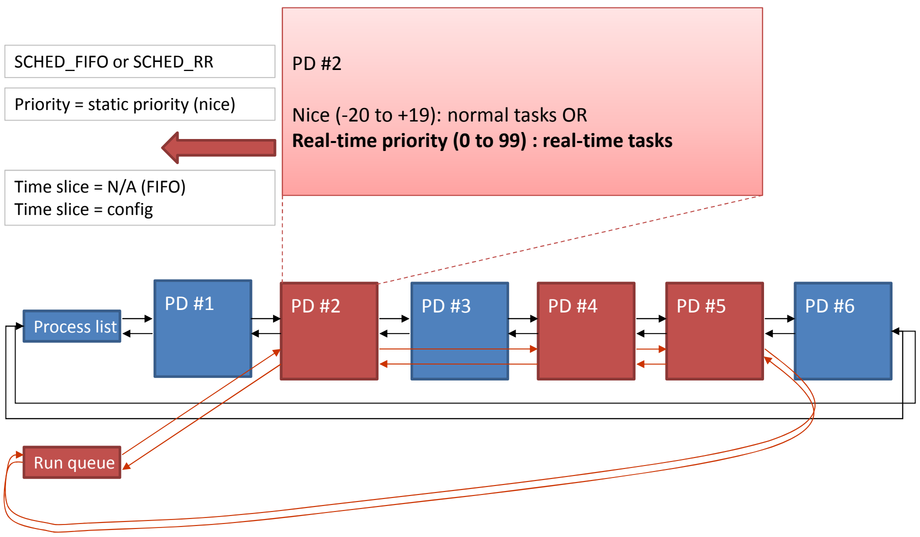

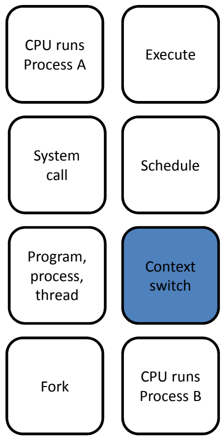


# Process schedule and context switching in Linux

- Context switch
  - Hardware context switch
    - Task State Segment Descriptor (Old Linux)
  - Step by step context switch
    - Better control and optimize
- Context switch
  - switch_mm()
    - Switch virtual memory mapping
  - switch_to()
    - Switch processor state
- Process switching occurs only in kernel mode
- The contents of all registers used by a process in User Mode have already been saved

# Scheduler in Details

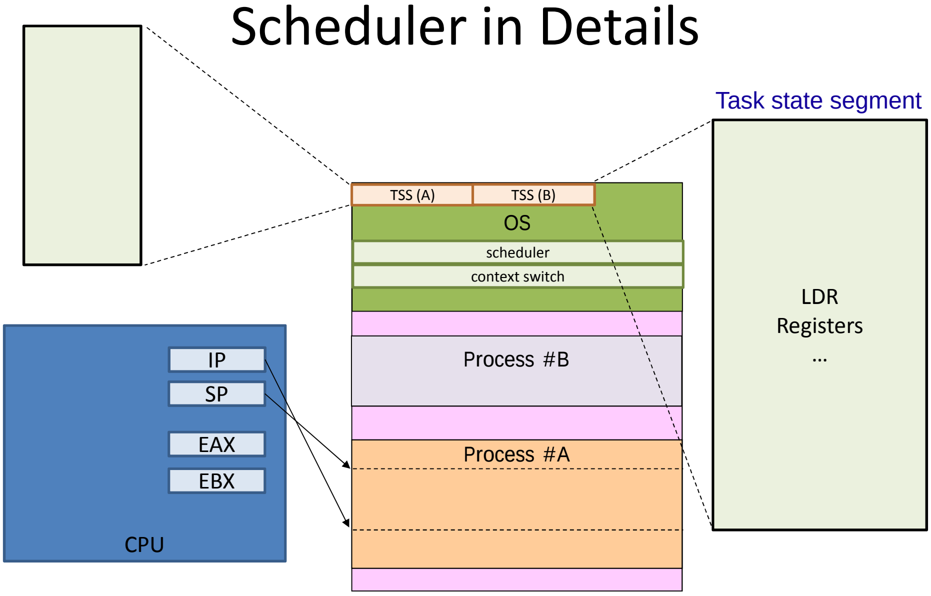


***TSS structure source: http://www.embedded.com/design/prototyping-and-development/4025054/Managing-Tasks-on-x86-Processors***

# Scheduler in Details

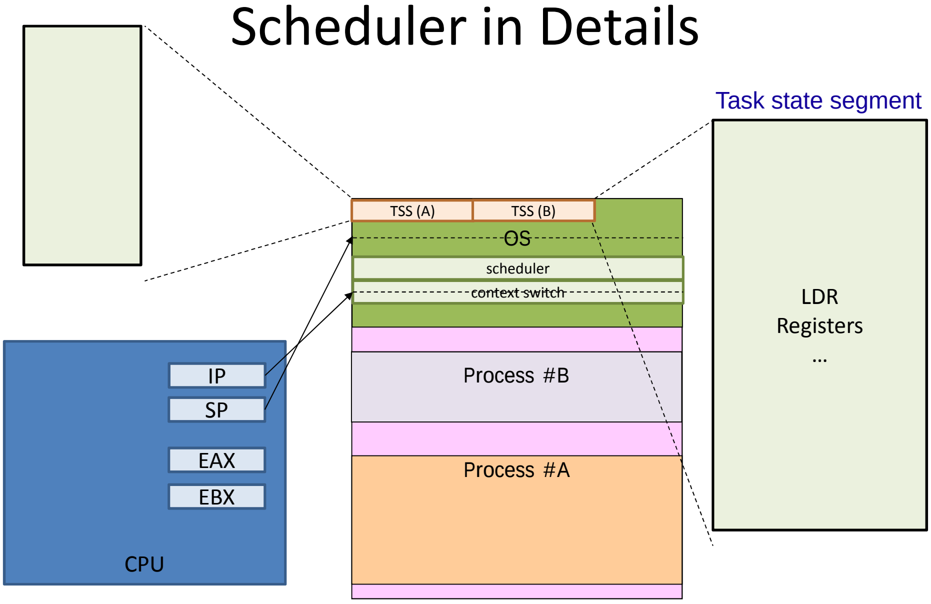


***TSS structure source: http://www.embedded.com/design/prototyping-and-development/4025054/Managing-Tasks-on-x86-Processors***

# Scheduler in Details

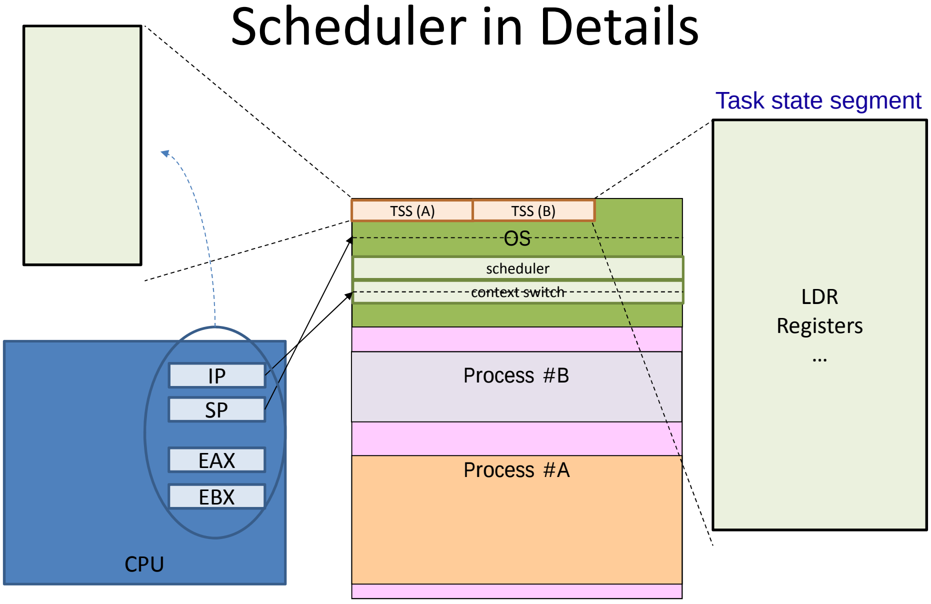


*TSS structure source: http://www.embedded.com/design/prototyping-and-development/4025054/Managing-Tasks-on-x86-Processors*

# Scheduler in Details

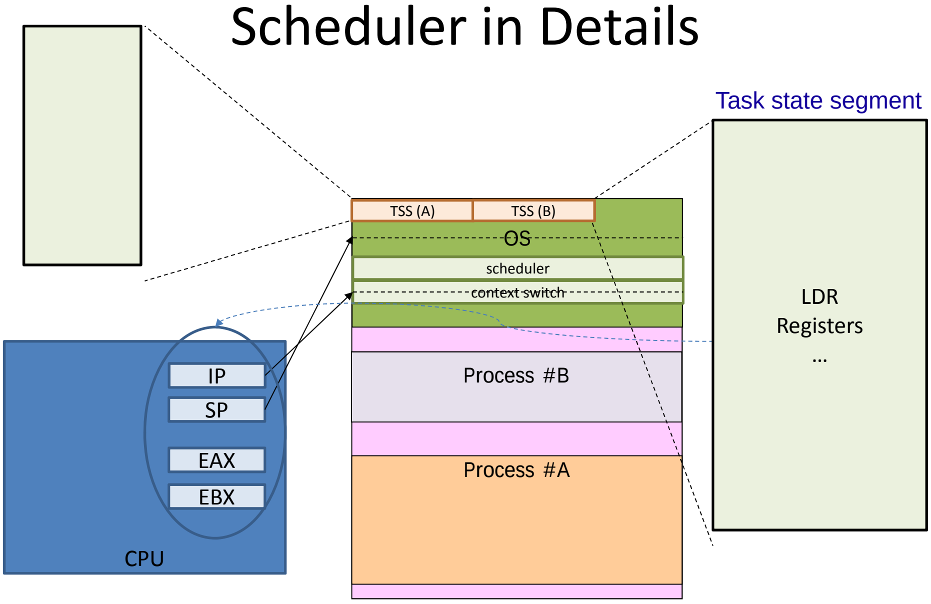


*TSS structure source: http://www.embedded.com/design/prototyping-and-development/4025054/Managing-Tasks-on-x86-Processors*

## Scheduler in Details

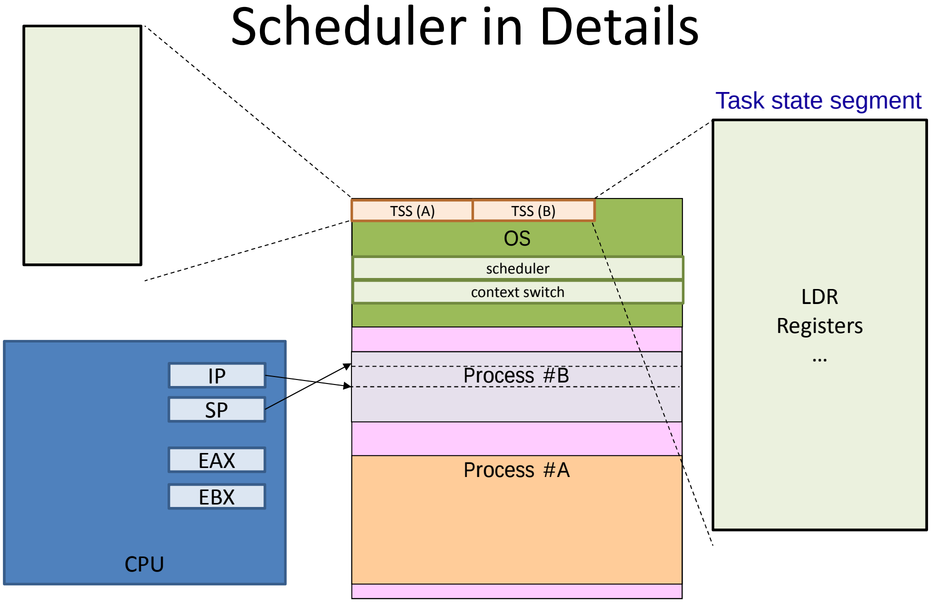


*TSS structure source: http://www.embedded.com/design/prototyping-and-development/4025054/Managing-Tasks-on-x86-Processors*

## How x86 helps in Context Switch

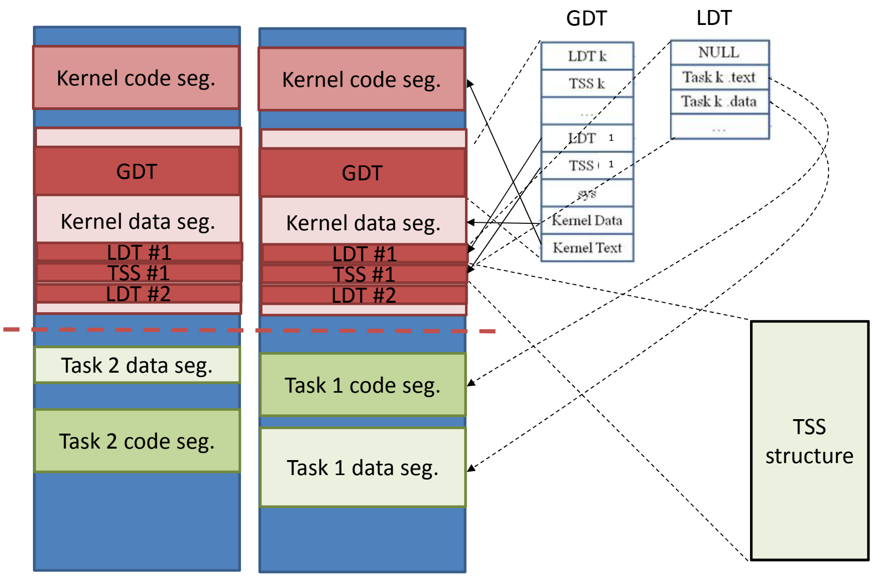

# Process related terms

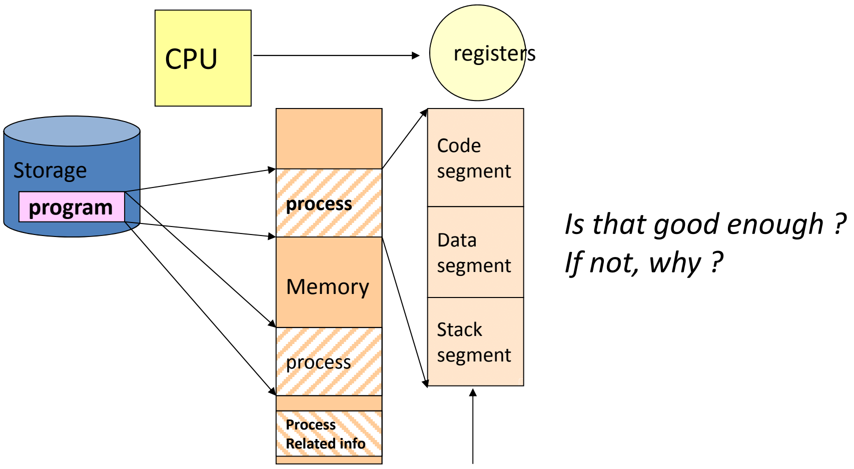


*Is that good enough ?*
*If not, why ?*

Physical memory might be discontinuous

47

# Process related terms (Cont.)

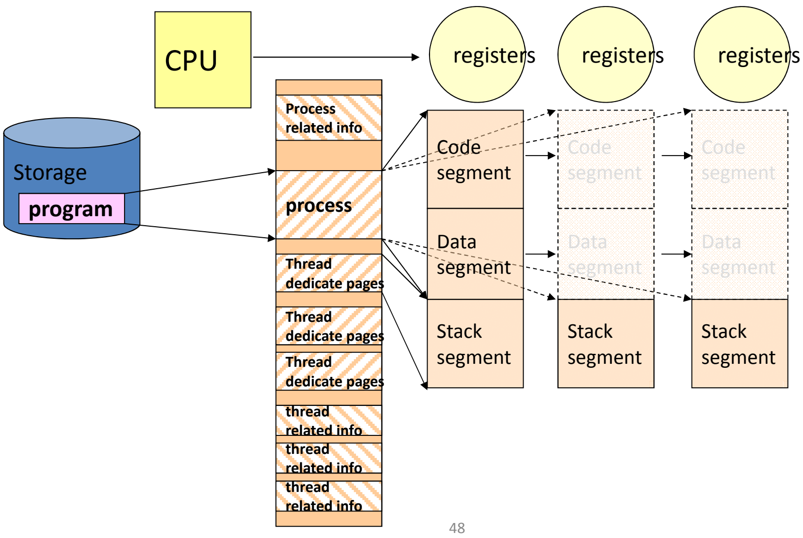


48

# Process related terms (Cont.)

- Depending on OS designs

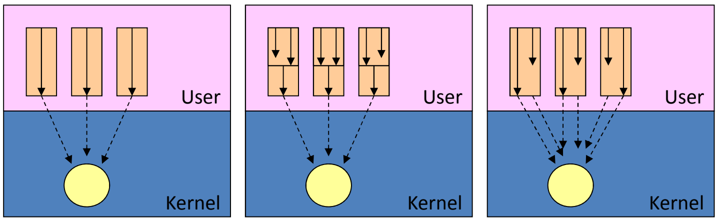


Processes only User thread kernel thread

49

# Process related terms (Cont.)

- Linux lightweight process

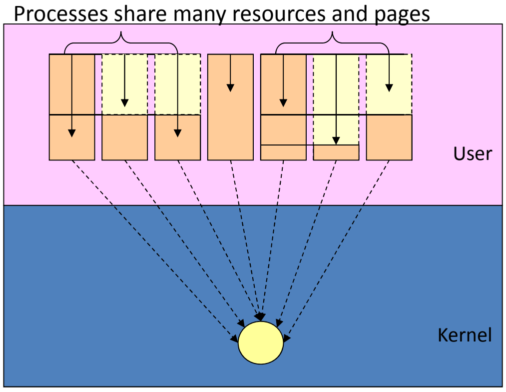


50