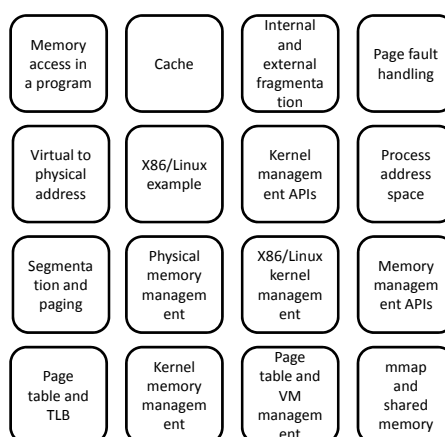
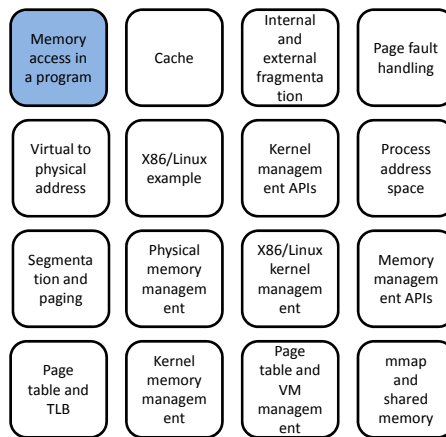


Operating System Design and Implementation

Memory Management – Part I

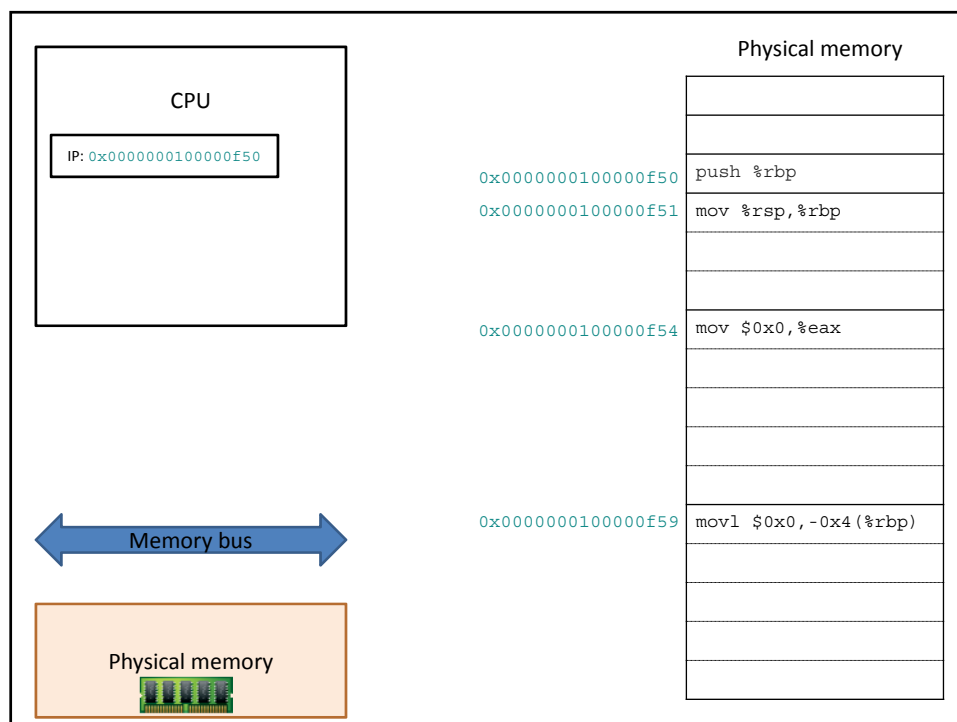
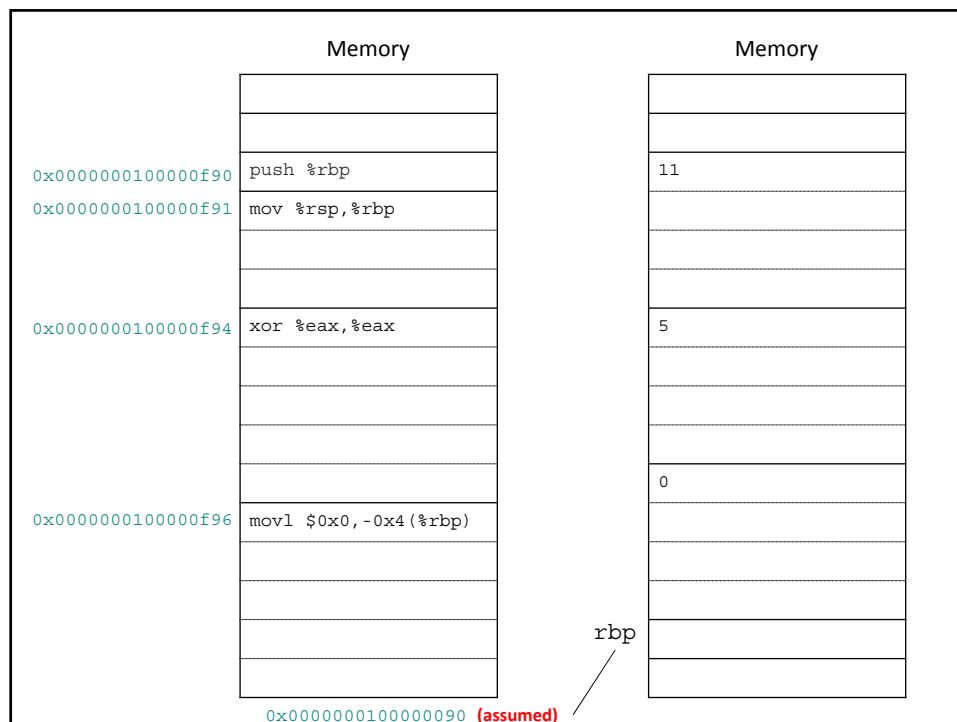
Shiao-Li Tsao





```
int main()
{
    int a = 5;
    int b = a + 6;
    return 0;
}
```

```
00000000100000f90 <_main>:
100000f90: 55                push    %rbp
100000f91: 48 89 e5          mov     %rsp,%rbp
100000f94: 31 c0            xor     %eax,%eax
100000f96: c7 45 fc 00 00 00 movl    $0x0,-0x4(%rbp)
100000f9d: c7 45 f8 05 00 00 movl    $0x5,-0x8(%rbp)
100000fa4: 8b 4d f8          mov     -0x8(%rbp),%ecx
100000fa7: 83 c1 06          add     $0x6,%ecx
100000faa: 89 4d f4          mov     %ecx,-0xc(%rbp)
100000fad: 5d                pop     %rbp
100000fae: c3                retq
```

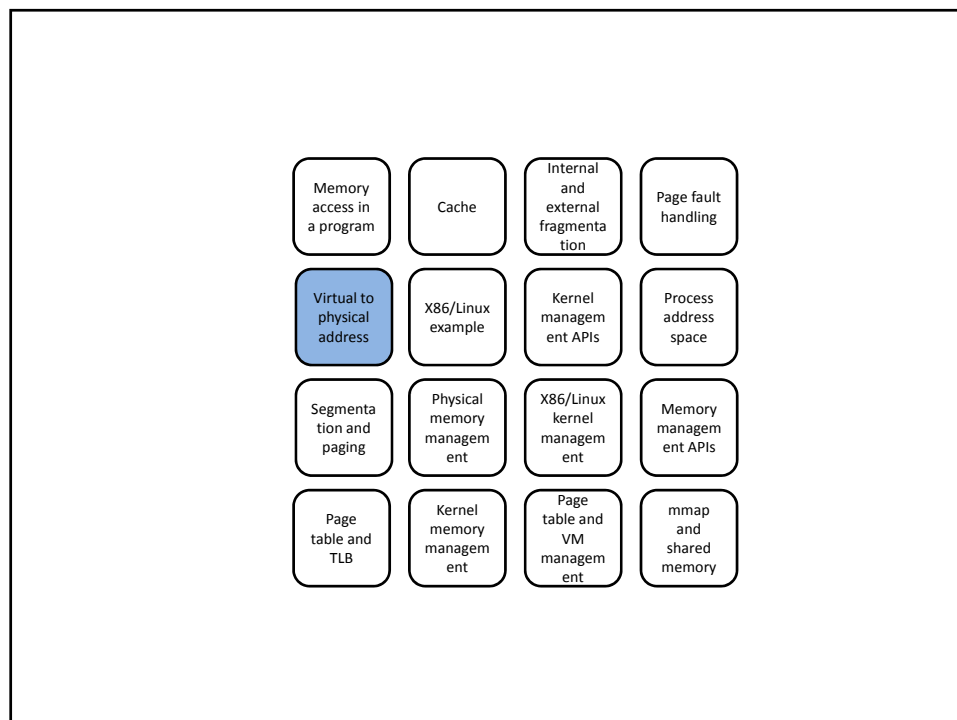
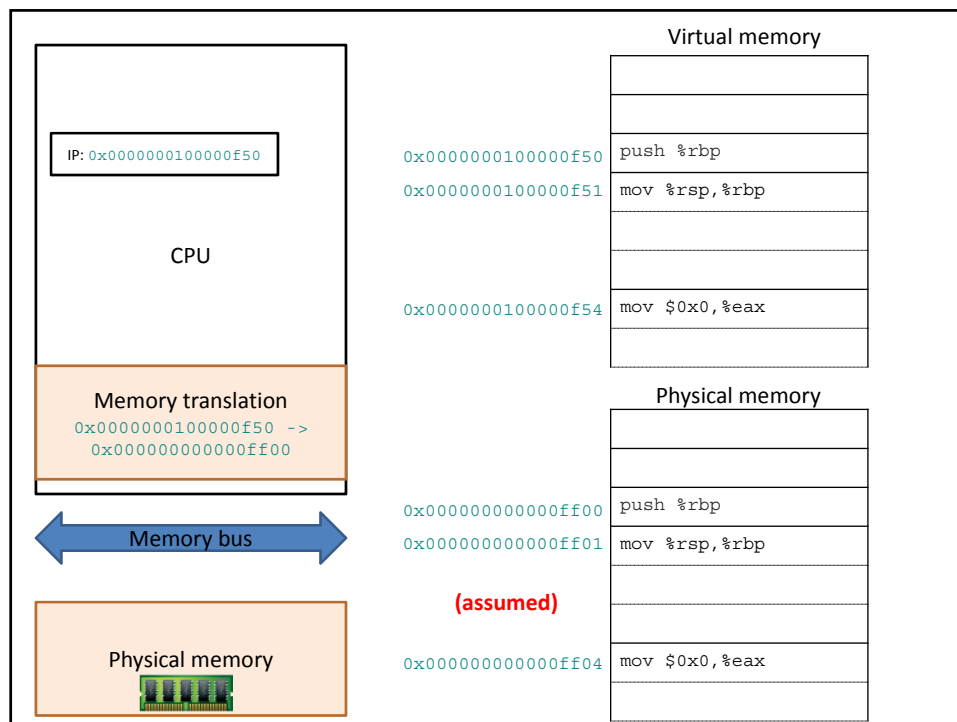


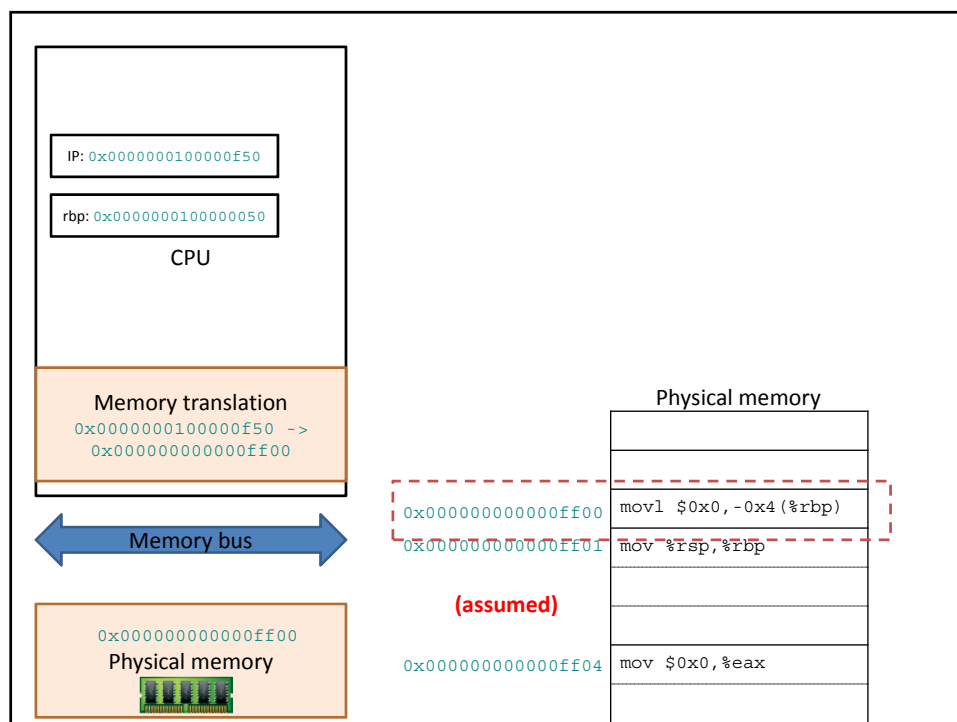
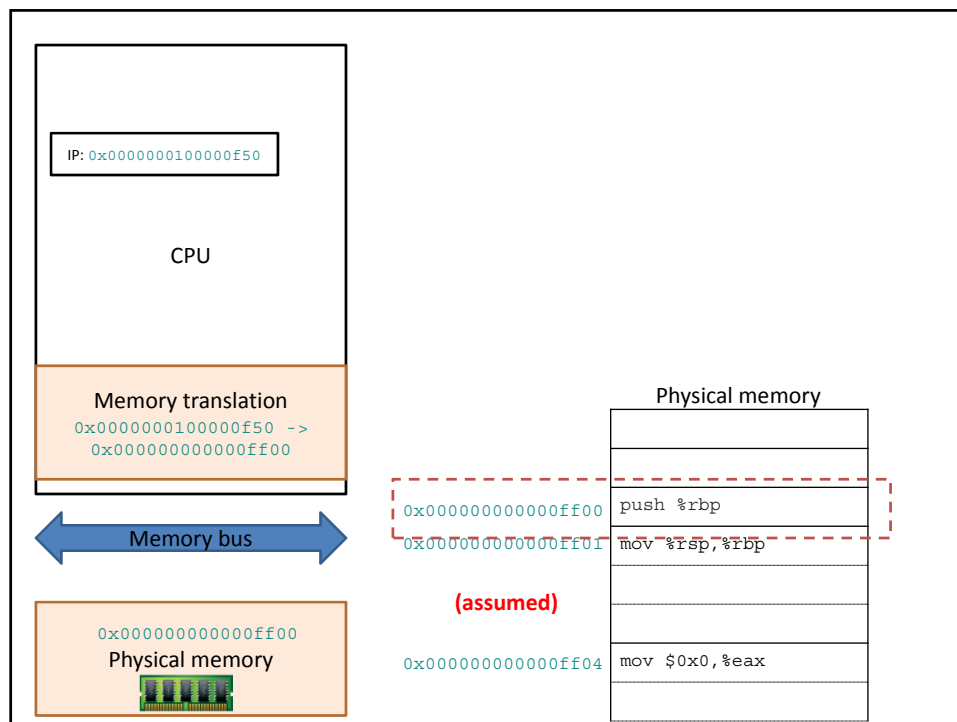
Physical Memory vs. Virtual Memory

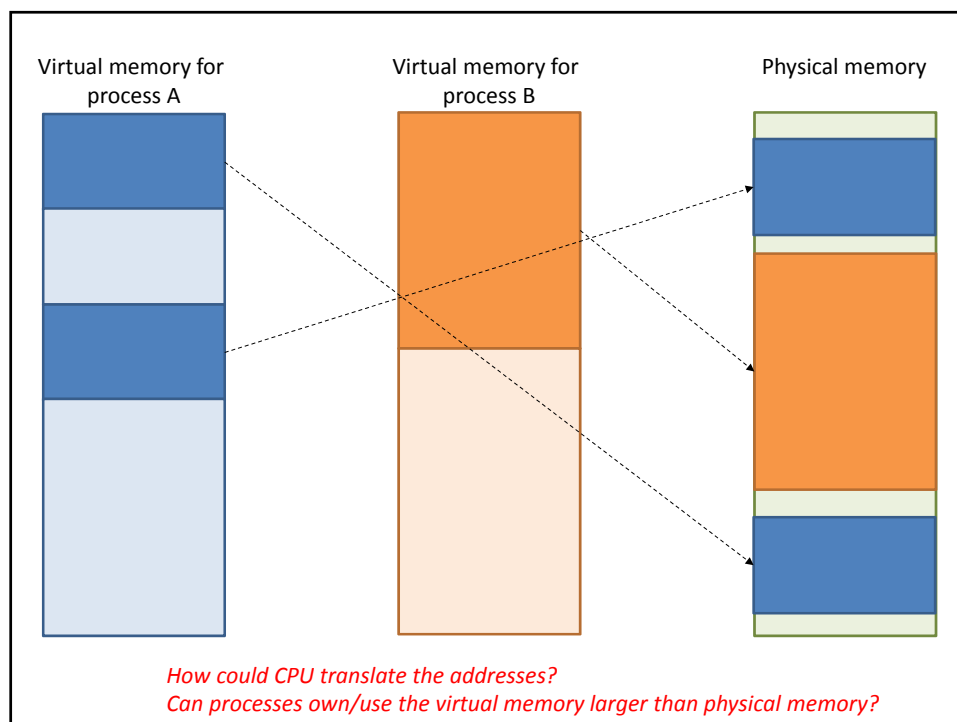
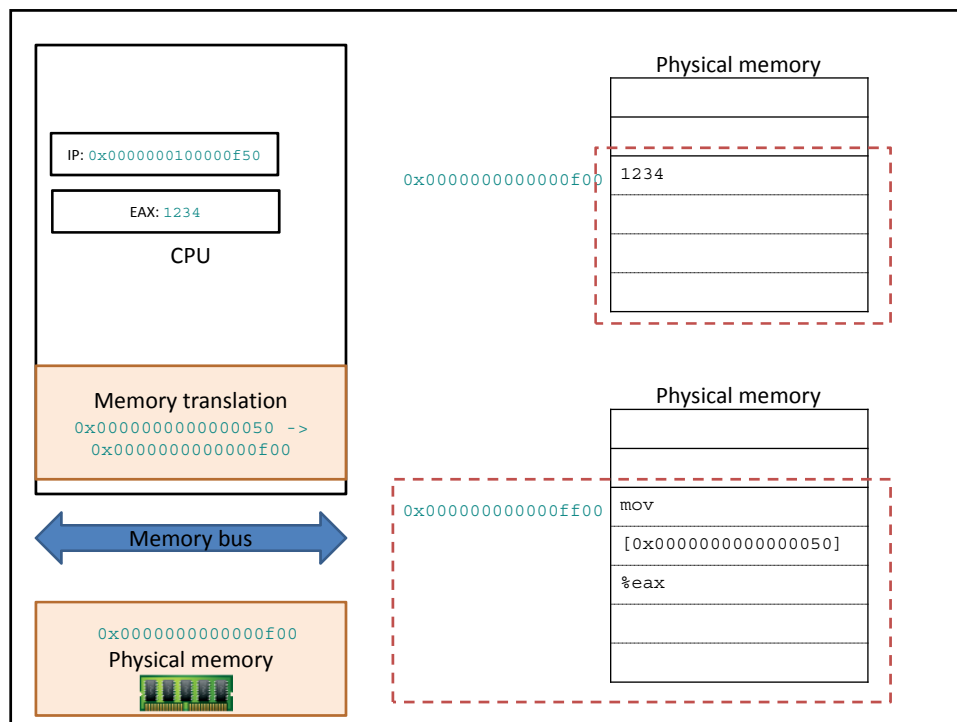
- Why not physical memory?
- Benefits for using virtual memory
 - Physical memory shared by multiple processes
 - Isolate processes from each other
 - Large/contiguous per process address space
 - Use memory more efficiently
- Disadvantage for using virtual memory
 - Expensive and slow
- *What kind of memory will you use if we only have single process, or you are the only developer of all processes, or you need very fast memory accesses*

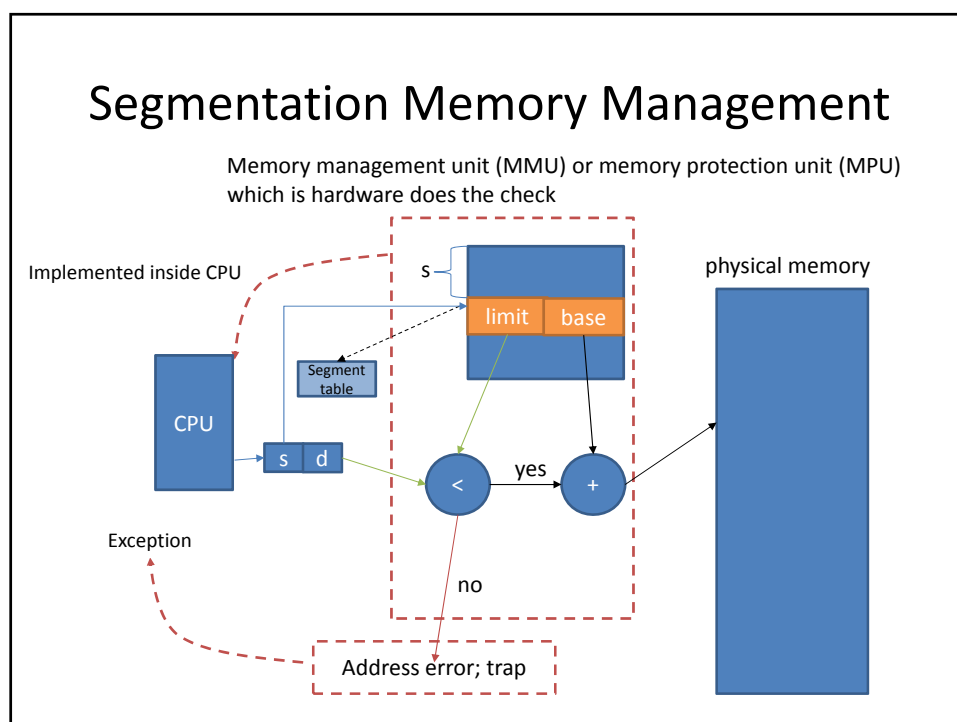
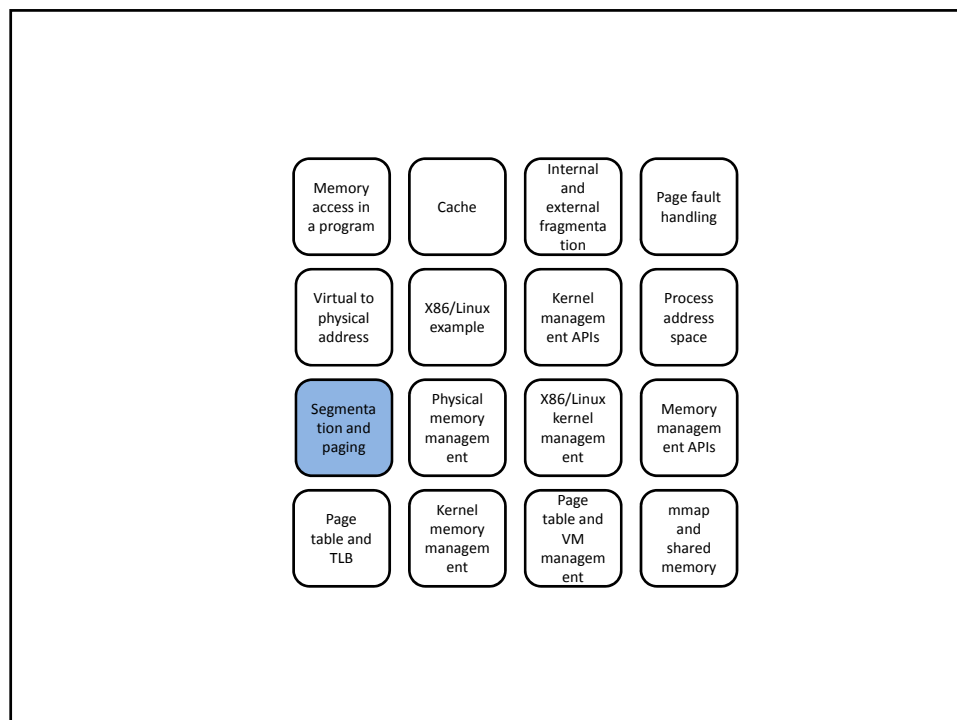
Virtual addresses

- The memory addresses users and compiler believed
- The code/data are stored in physical memory
- The virtual memory addresses are translated into physical memory addresses (by CPU) and then CPU can access the code/data
- How can a CPU translate the addresses?
 - Offset? Table lookup?








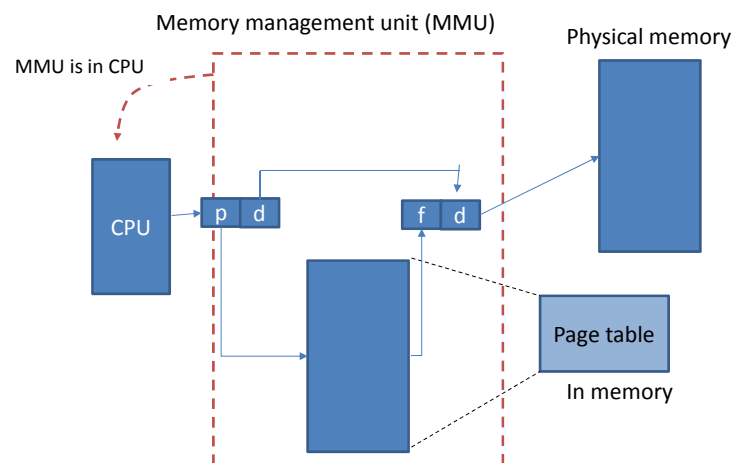


Segmentation Memory Management (example)

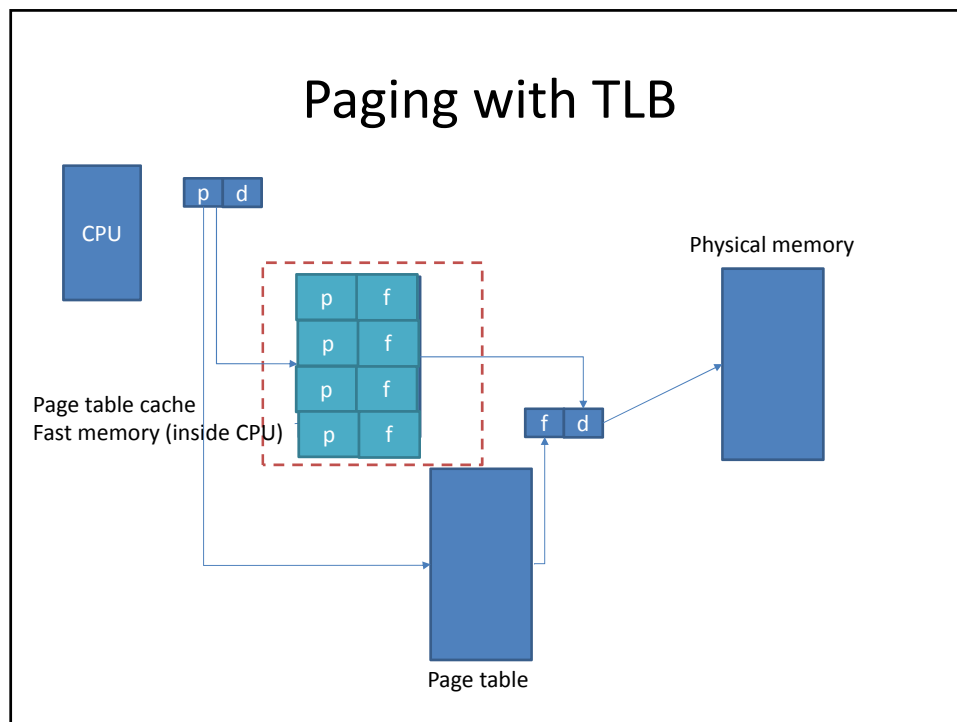
segment	limit	base
3	0x2000	0x1234
1	0x1234	0x1200



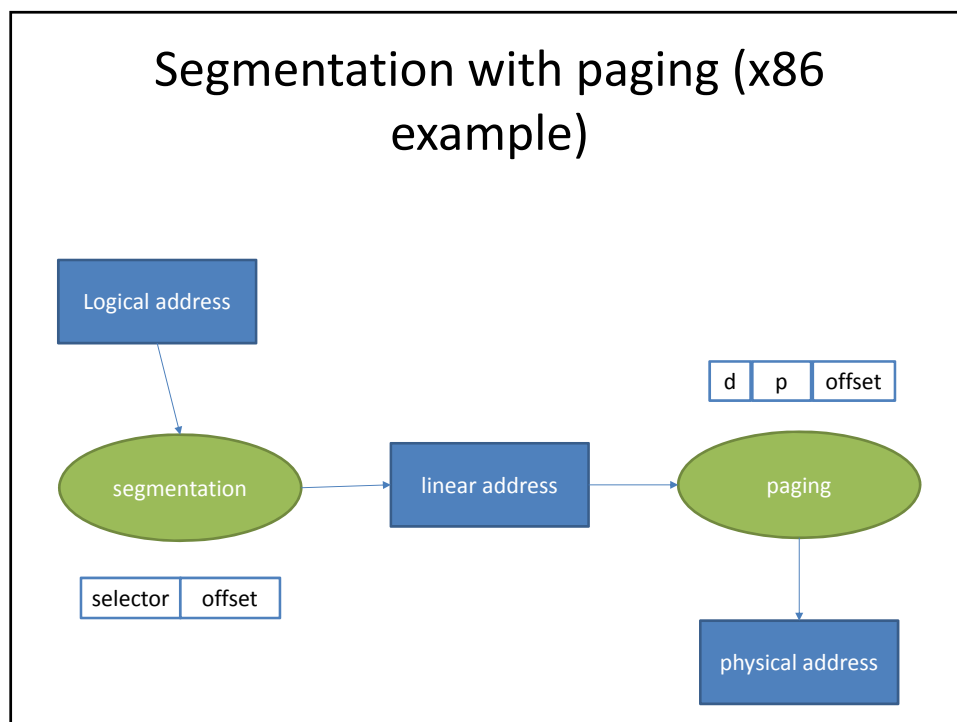
Page memory management

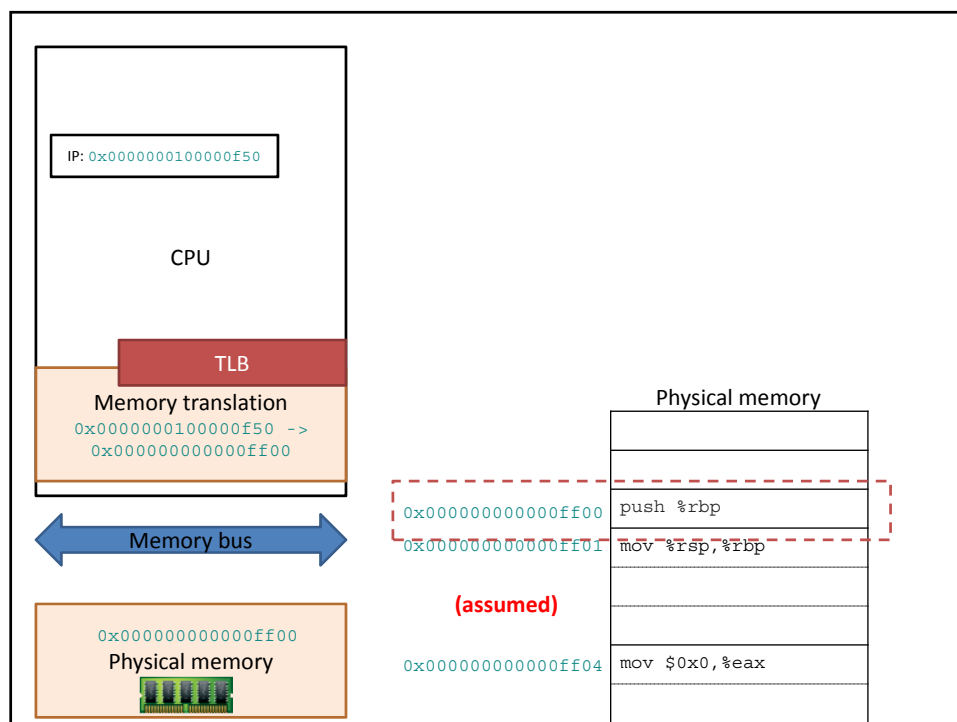
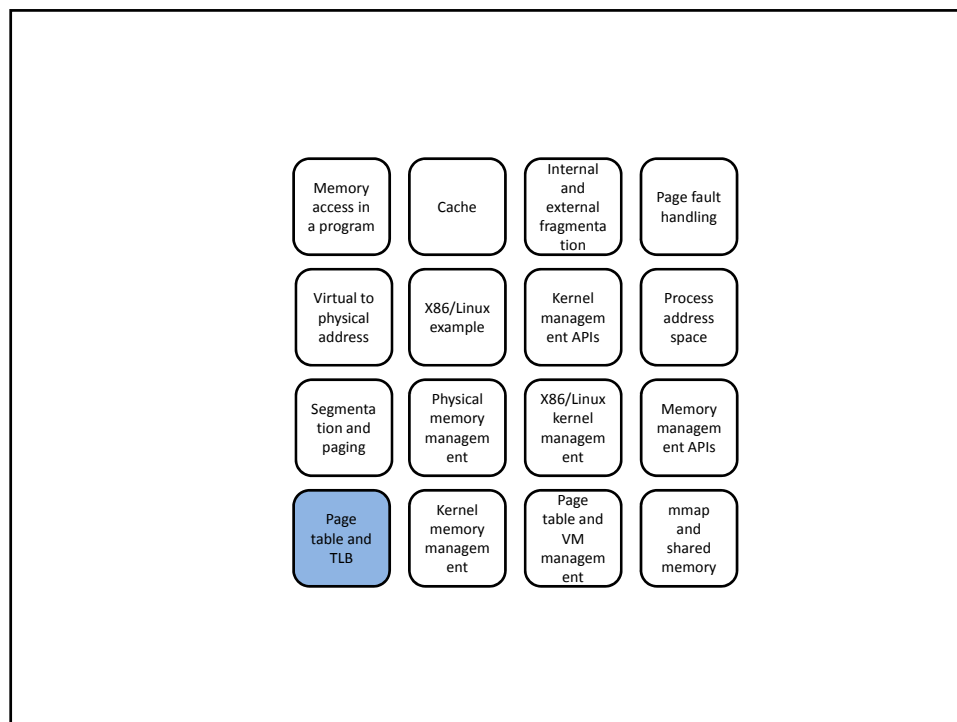


Paging with TLB

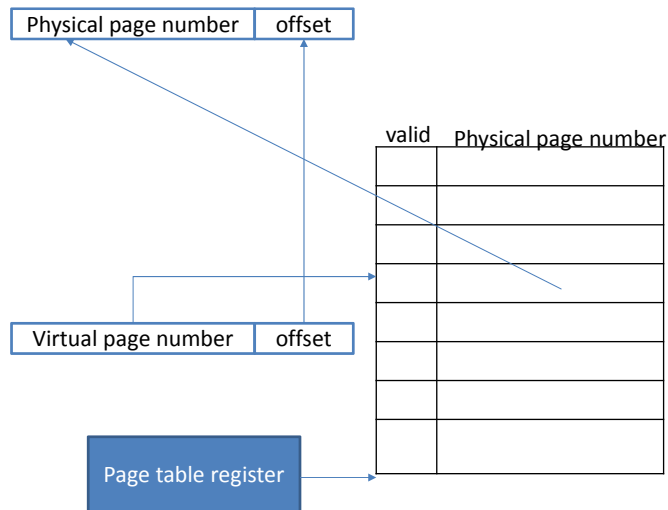


Segmentation with paging (x86 example)

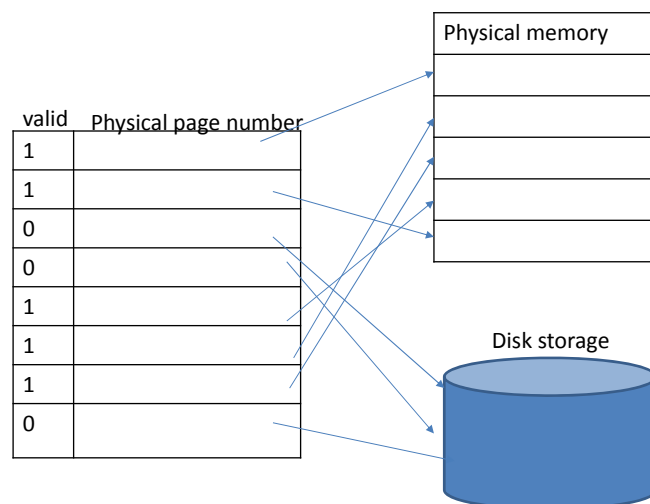


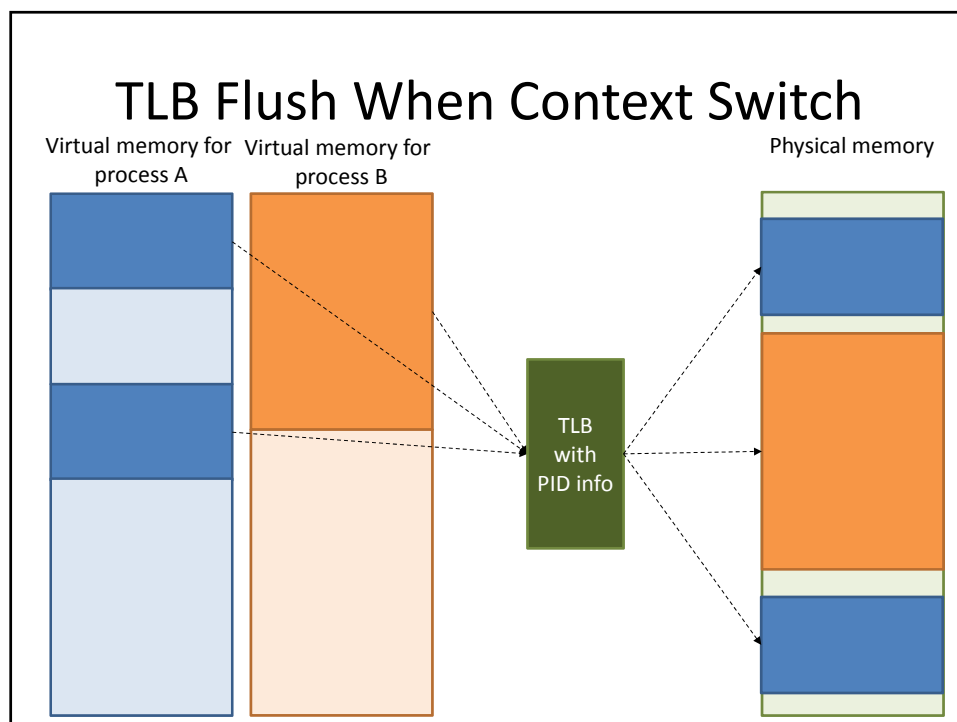
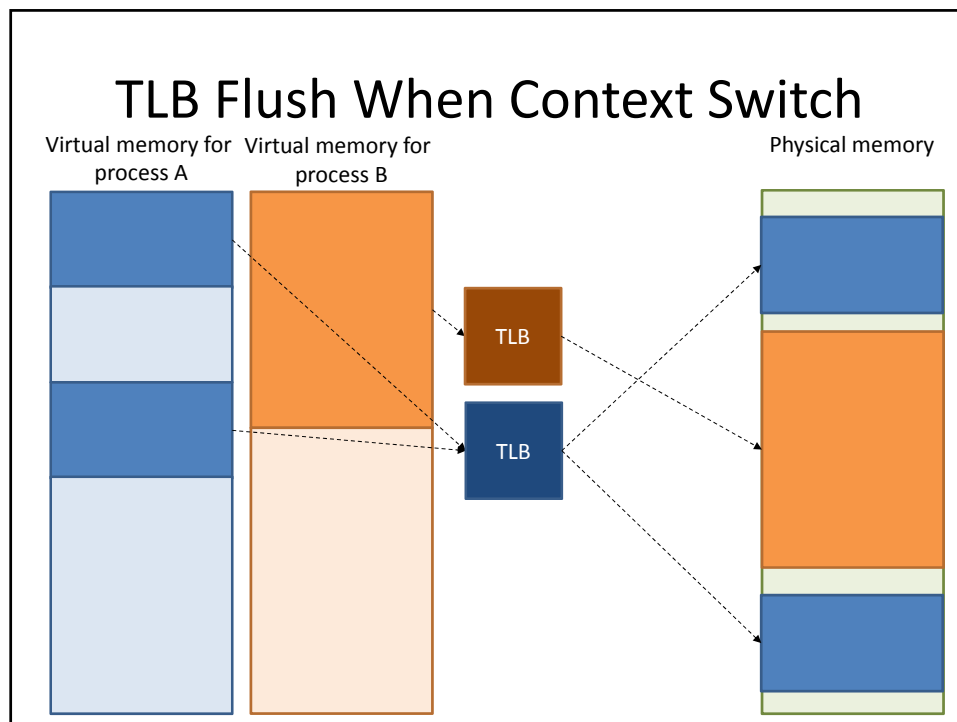


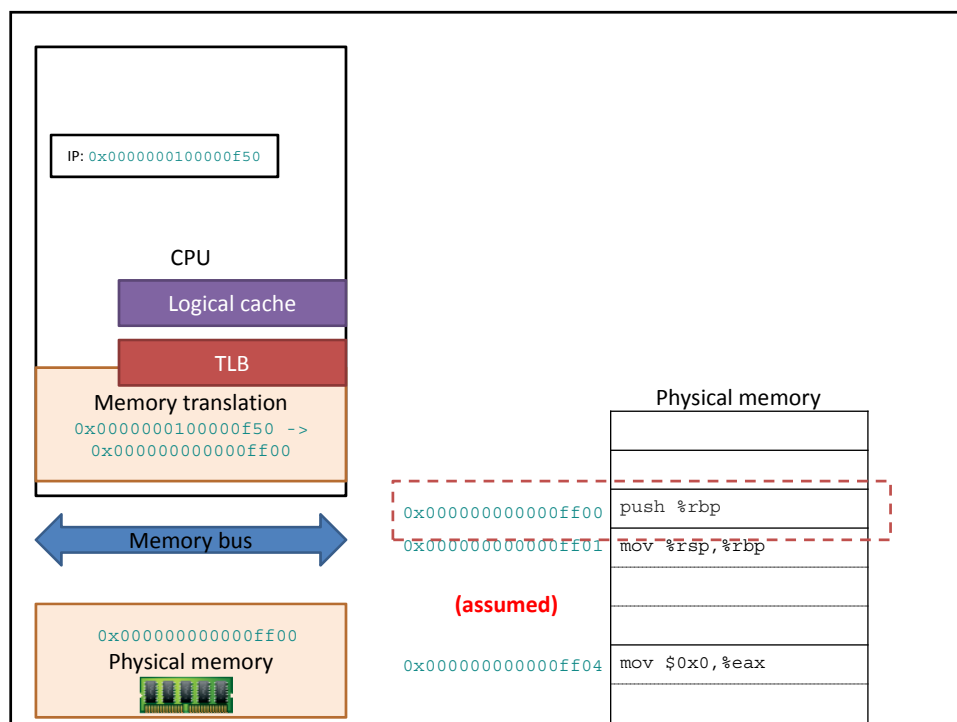
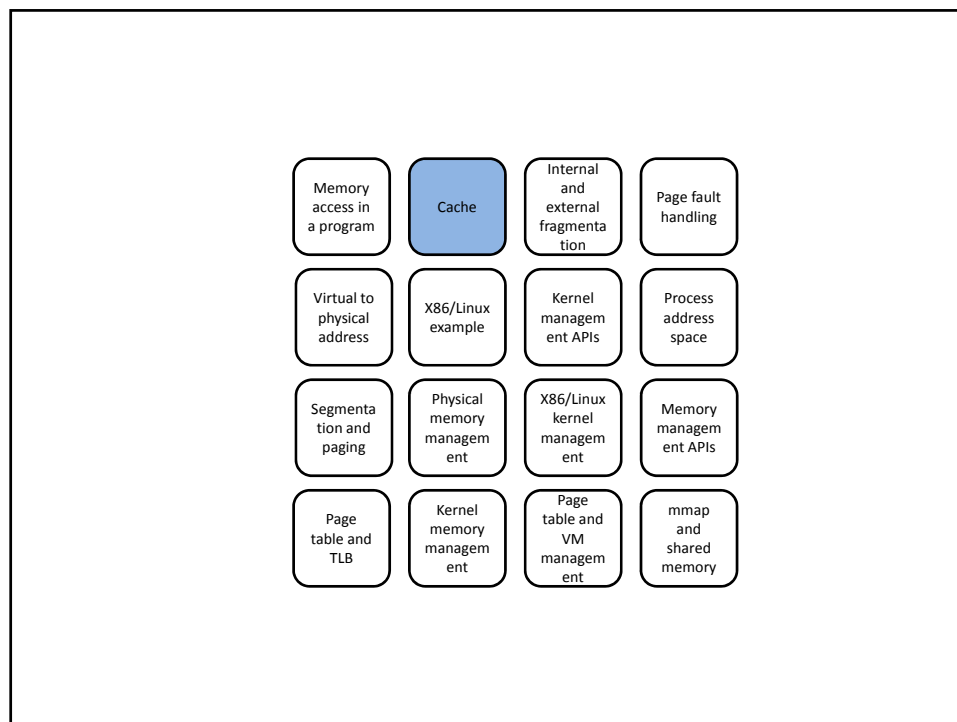
Page Table

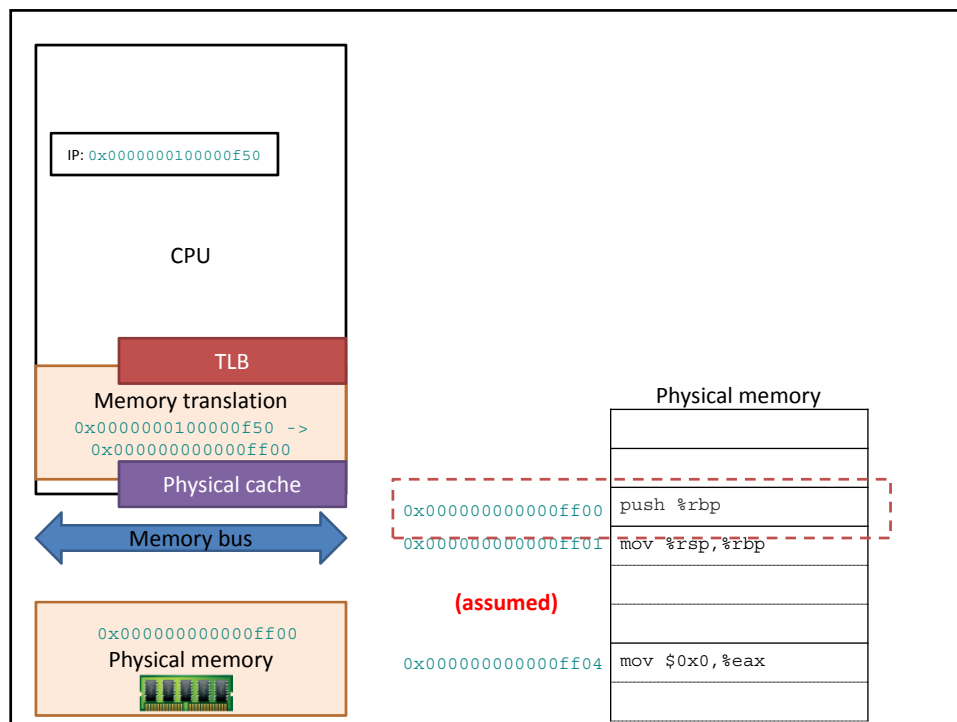


Page Table and Virtual Address

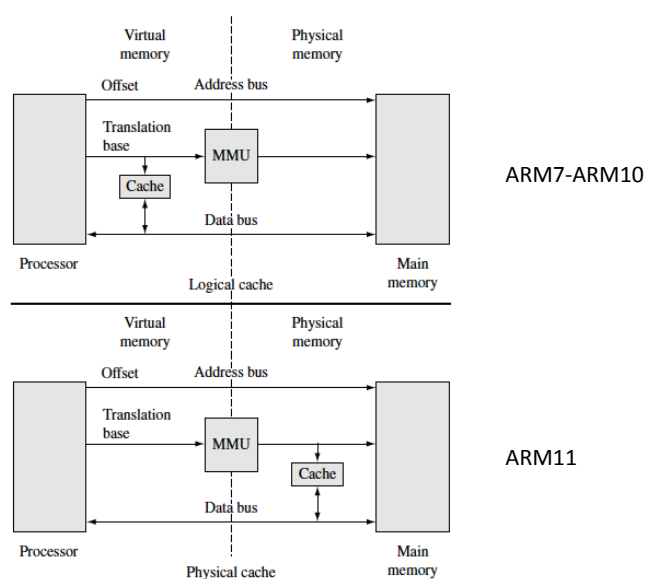








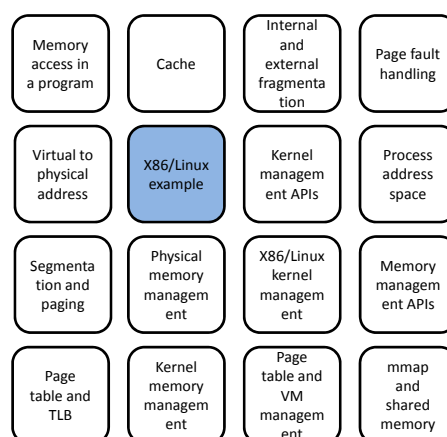
Logical and Physical Memory


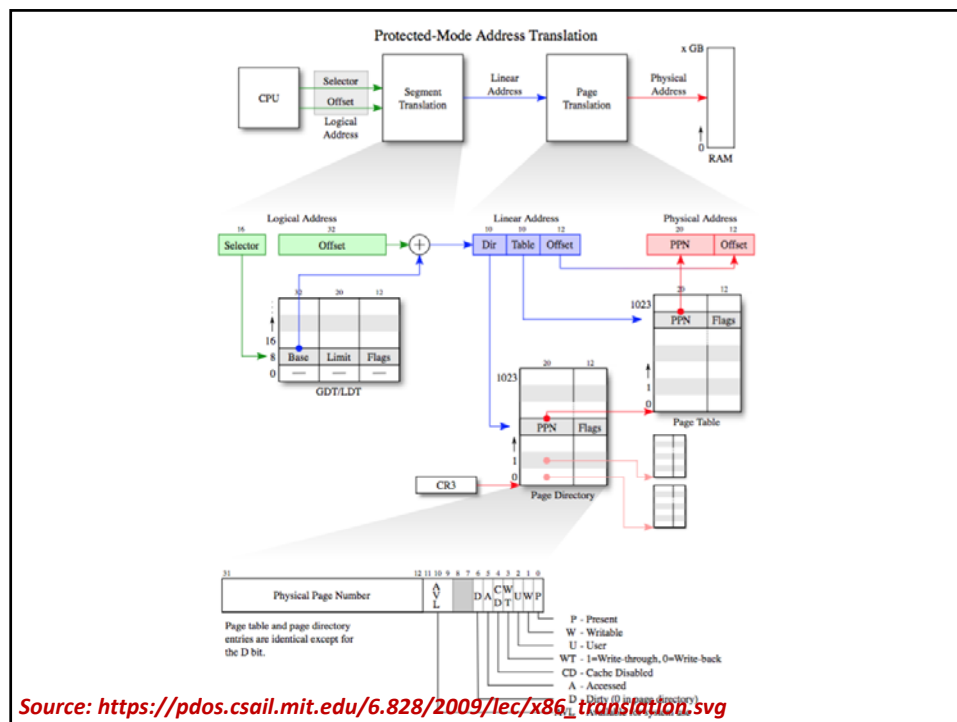


Logical and Physical Memory

- Physical Caches
 - Slow but share without flush
- Logical Caches
 - Pid or flush while context switch

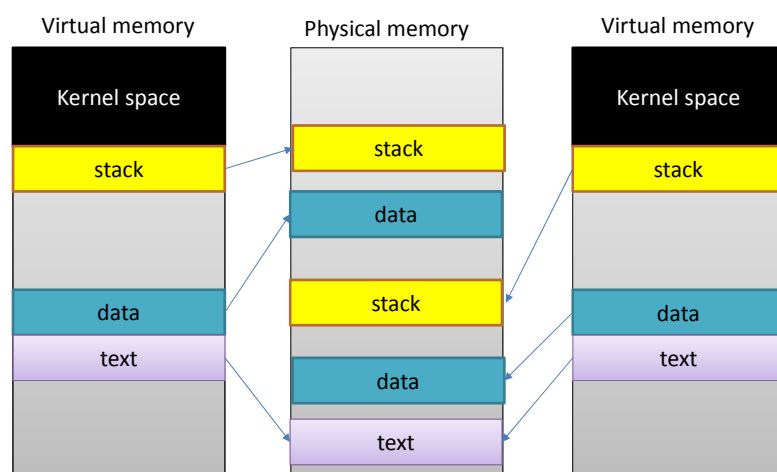
31



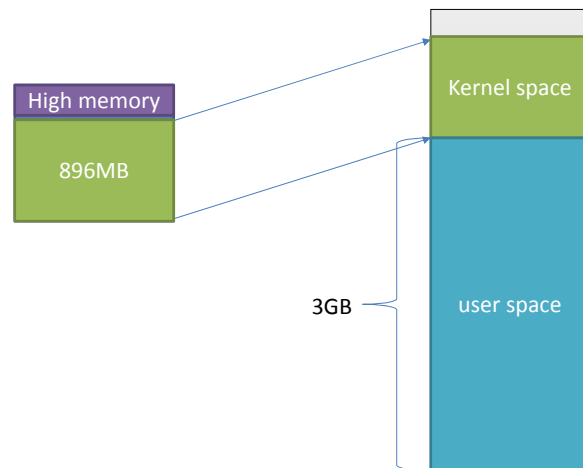


User Address Space

- 4G for 32 bits processor
- How about kernel?
 - Kernel uses another 4G address space ?
 - Context switch for system calls?
- How about user and kernel share the same 4G spaces?
 - User program can use X
 - Kernel program can use $4G-X$
- Can user access the kernel data?
 - Limited by user data segment
- Can kernel access the user data?
 - Kernel data segment covers 4G

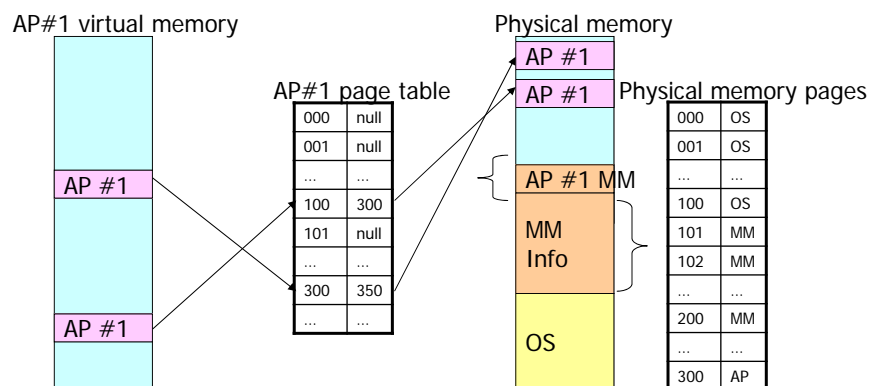


Direct linear mapping



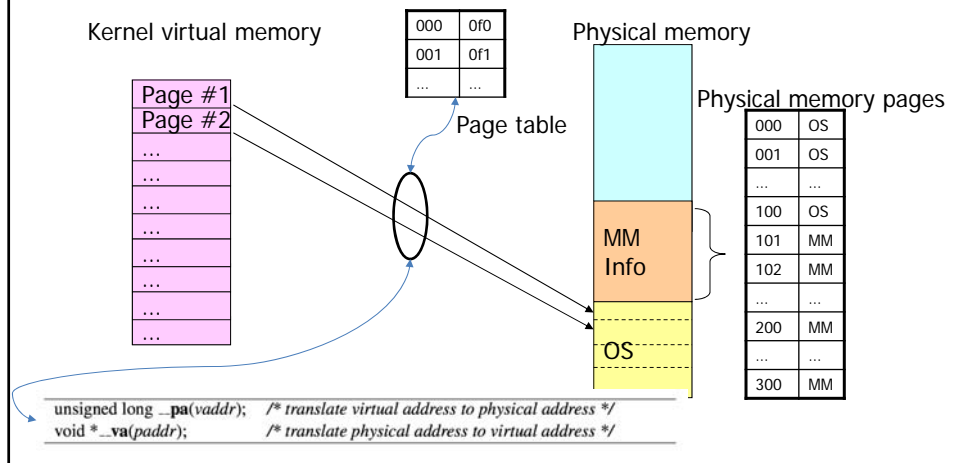
MM Basics (Cont.)

- Management physical memory (page)



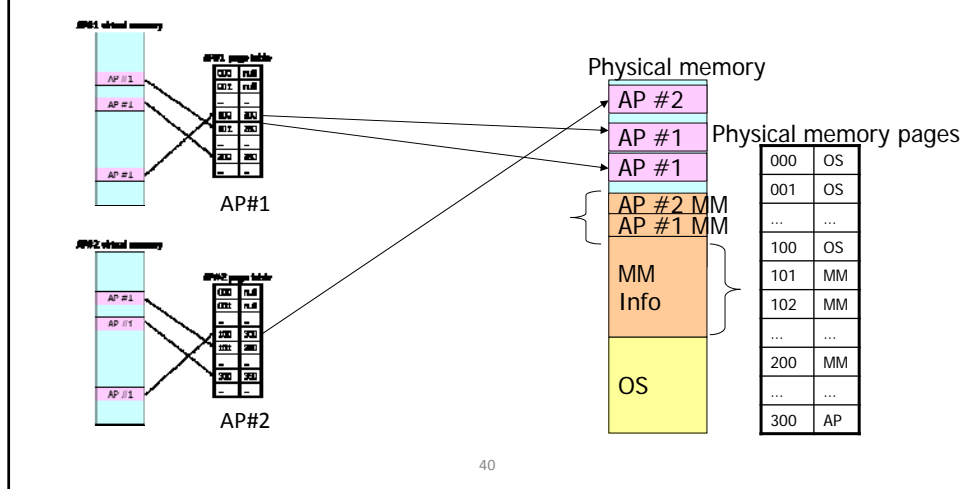
MM Basics (Cont.)

- Kernel memory allocation/release



MM Basics (Cont.)

- application memory allocation/release



40

