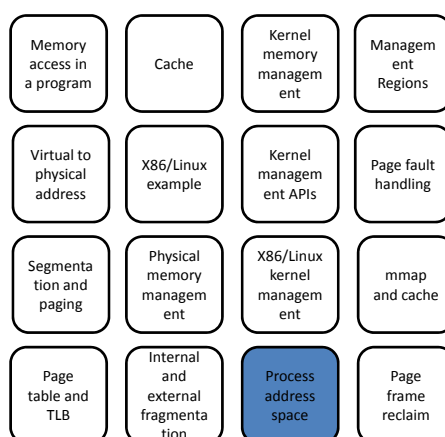


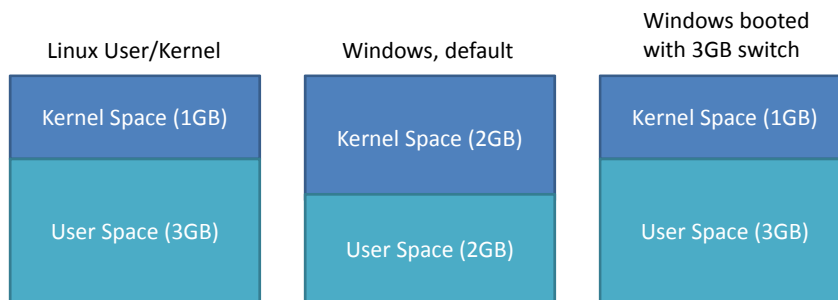
Operating System Design and Implementation

Memory Management – Part III


Shiao-Li Tsao



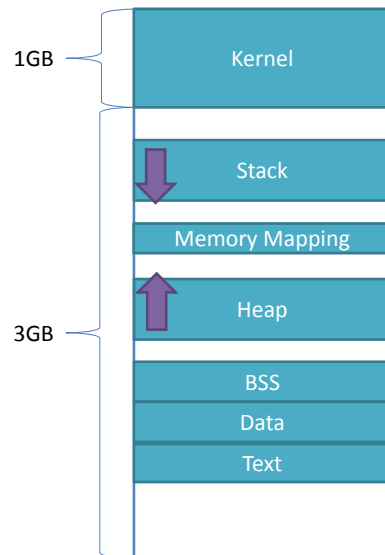
Kernel Space vs. User Space




Virtual Addresses Before/After Context Switches


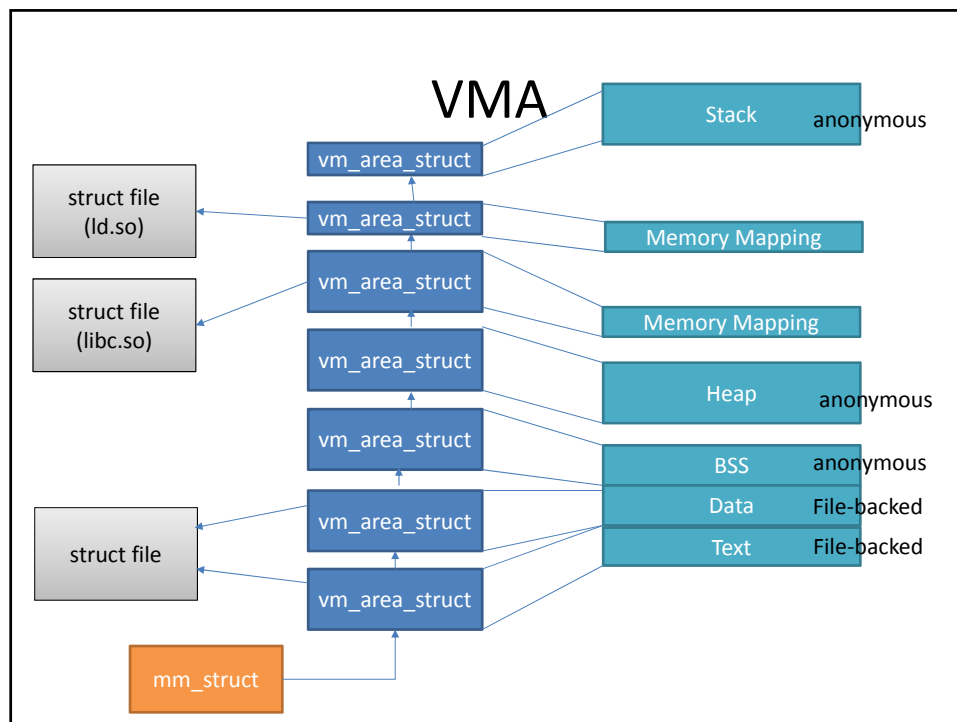


Per Process Address Space

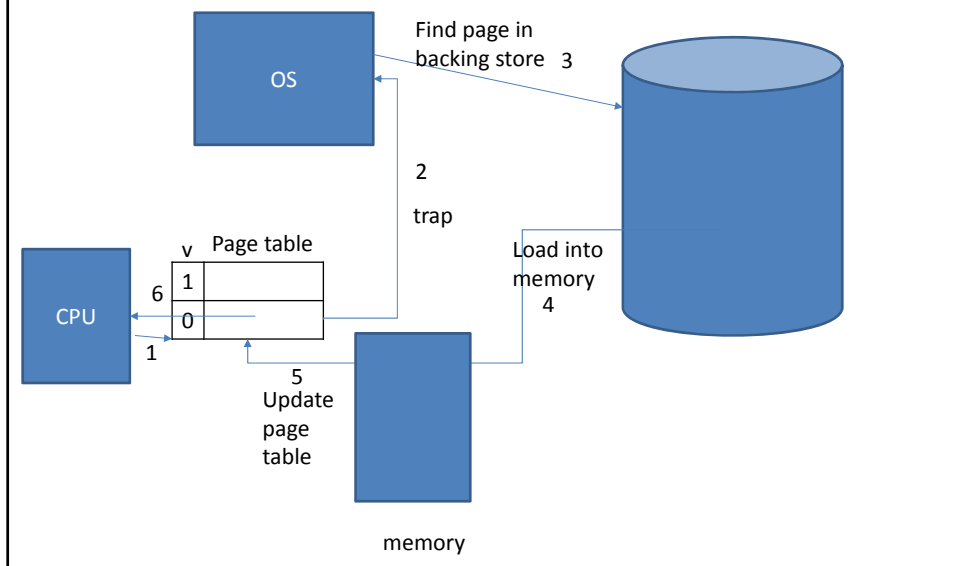


task_struct and mm_struct





Page Fault

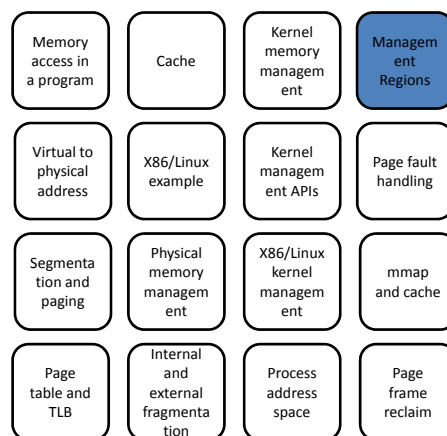


Allocation of MM Related Structures


```
slabinfo - version: 1.1 (SMP)
kmem_cache      80      80      248      5      5      1 : 252 126
urb_priv         0       0       64      0      0      1 : 252 126
tcp_bind_bucket 15      226      32      2      2      1 : 252 126
inode_cache     5714    5992     512    856    856      1 : 124  62
dentry_cache    5160    5160     128    172    172      1 : 252 126
mm_struct        240     240     160     10     10      1 : 252 126
vm_area_struct  3911    4480      96    112    112      1 : 252 126
size-64(DMA)     0       0       64      0      0      1 : 252 126
size-64          432    1357      64     23     23      1 : 252 126
size-32(DMA)     17     113      32      1      1      1 : 252 126
size-32          850    2712      32     24     24      1 : 252 126
```

Linux Kernel Thread and mm_struct

- Kernel threads do not have a process address space
 - mm field of a kernel thread's process descriptor is NULL
- Lack of an address space is fine, because kernel threads do not ever access any user-space memory
- Better performance

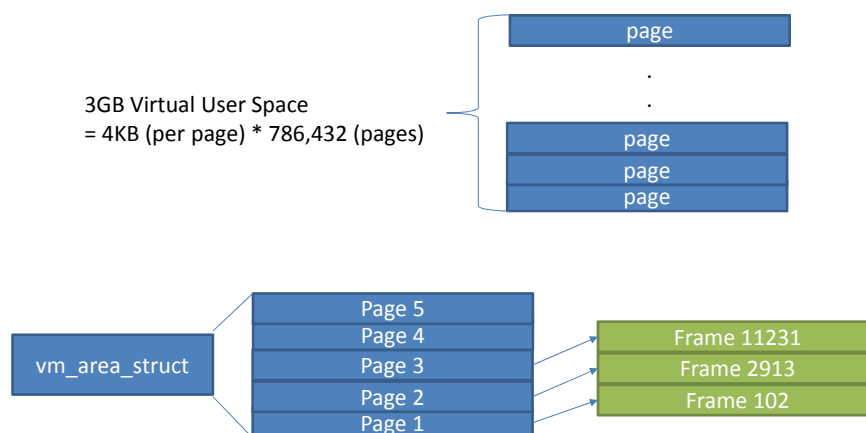


mm_struct, vm_area_struct, and page




VMA, Page and Frame

3GB Virtual User Space
= 4KB (per page) * 786,432 (pages)




Memory Region




Red-Black Tree

- $O(\log n)$



By Cburnett,CC-BY-SA-3.0, via Wikimedia Commons"

VMA Search Using R-B Tree

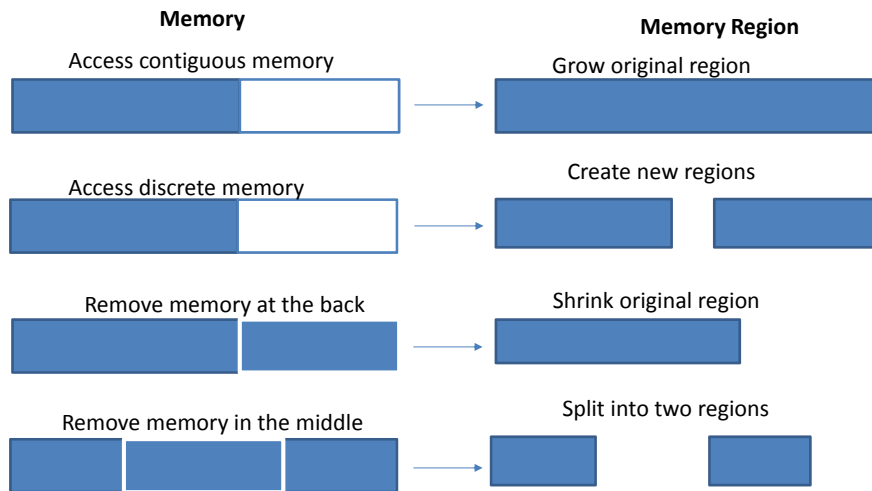


Ref: <http://www.cs.columbia.edu/~junfeng/13fa-w4118/lectures/l20-adv-mm.pdf>

malloc() and Heap Management




Process Address Space

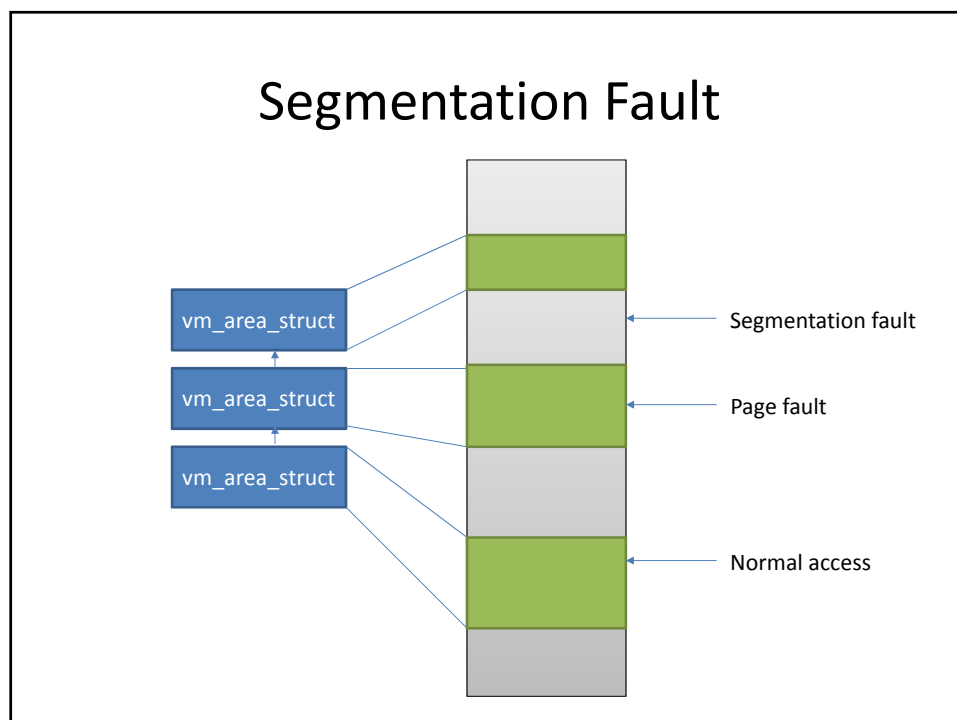
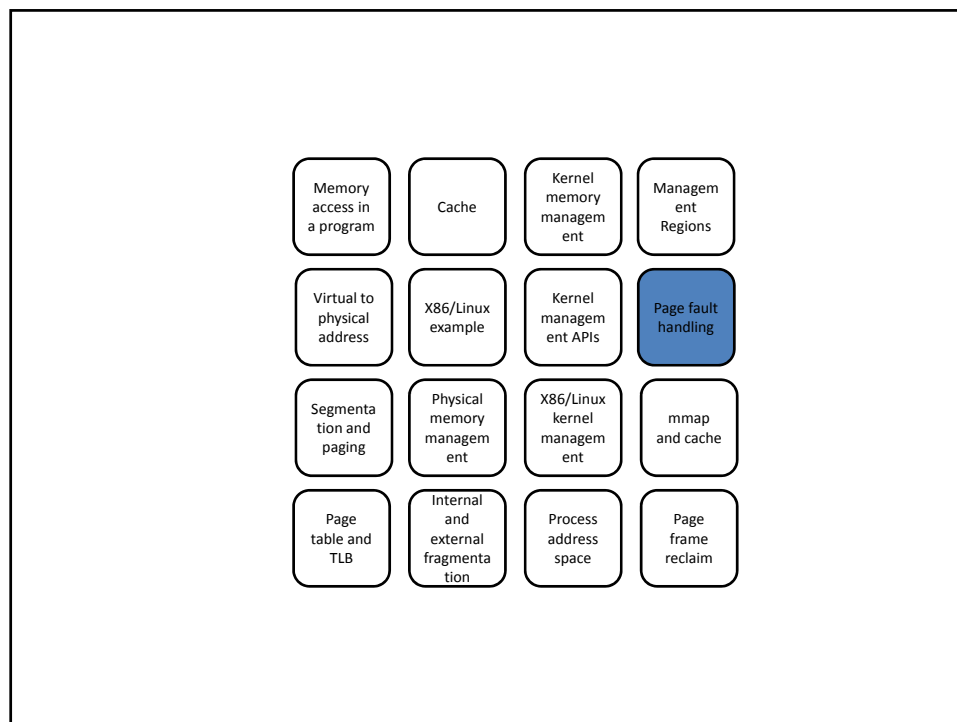


19

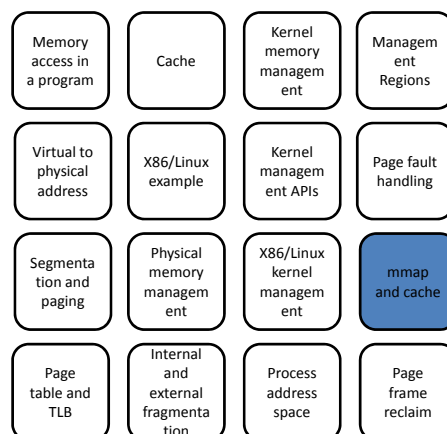
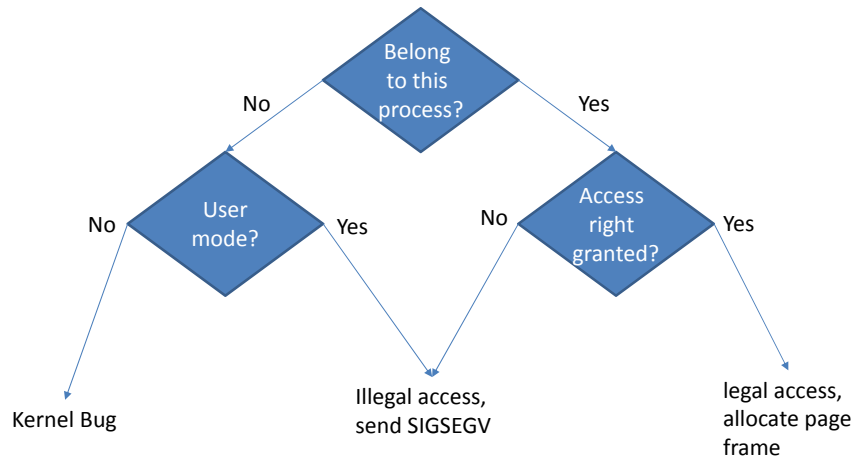
Enlarge VMA




Ref: <http://duartes.org/gustavo/blog/post/how-the-kernel-manages-your-memory/>




Page Fault Handler



Memory Mapped



Data Structure for File Memory Mapping



Ref: Understanding the Linux Kernel, Daniel P. Bovet, Marco Cesati

Two Types of Memory Mapping

- A file mapping maps a memory region to a region of a file
 - backing store = file
 - as long as the mapping is established, the content of the file can be read from or written to using direct memory access (“as if they were variables”)
- An anonymous mappings maps a memory region to a fresh “virtual” memory area filled with 0
 - backing store = zero-ed memory area

Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchioli

Having memory mapped pages in common

- Thanks to virtual memory management, different processes can have mapped pages in common
- More precisely, mapped pages in different processes can refer to physical memory pages that have the same backing store
- That can happen in two ways:
 - through fork, as memory mappings are inherited by children
 - when multiple processes map the same region of a file


Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchioli

Shared vs private mappings

- With mapped pages in common, the involved processes might see changes performed by others to mapped pages in common, depending on whether the mapping is:
 - private mapping in this case modifications are not visible to other processes.
 - pages are initially the same, but modification are not shared, as it happens with copy-on-write memory after fork
 - private mappings are also known as copy-on-write mappings
 - shared mapping in this case modifications to mapped pages in common are visible to all involved processes
 - i.e. pages are not copied-on-write

Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchioli

File mapping and memory layout



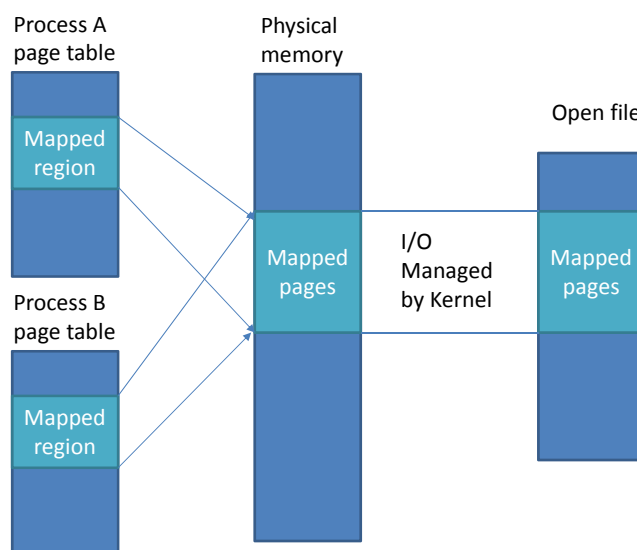
Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchioli

Shared file mapping (#2)

- Effects:
 - processes mapping the same region of a file share physical memory frames
 - more precisely: they have virtual memory pages that map to the same physical memory frames
 - additionally, the involved physical frames have the mapped file as ultimate backing store
 - i.e. modifications to the (shared) physical frames are saved to the mapped file on disk

Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchioli

Shared file mapping (#2)



Shared file mapping (#2) (cont.)

- Use cases
 - memory-mapped I/O, as an alternative to read/write
 - as in the case of private file mapping, but here it works for both reading and writing data
 - Inter-process communication, with the following characteristics:
 - data-transfer (not byte stream)
 - with filesystem persistence
 - among unrelated processes

Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchioli

Memory-mapped I/O

- Given that:
 - memory content is initialized from file
 - changes to memory are reflected to file
- we can perform I/O by simply changing bytes of memory.
- Access to file mappings is less intuitive than sequential read/write operations
 - the mental model is that of working on your data as a huge byte array (which is what memory is, after all)
 - a best practice to follow is that of defining struct-s that correspond to elements stored in the mapping, and copy them around with memcpy & co

Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchioli

Memory-mapped I/O — advantages

- performance gain: 1 memory copy
 - with read/write I/O each action involves 2 memory copies: 1 between user-space and kernel buffers + 1 between kernel
- buffers and the I/O device
 - with memory-mapped I/O only the 2nd copy remains
 - flash exercise: how many copies for standard I/O?
- performance gain: no context switch
 - no syscall and no context switch is involved in accessing mapped memory
 - page faults are possible, though reduced memory usage
 - we avoid user-space buffers ! less memory needed
 - if memory mapped region is shared, we use only one set of
- buffers for all processes seeking is simplified
 - no need of explicit lseek, just pointer manipulation


Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchiroli

Memory-mapped I/O — disadvantages

- memory garbage
 - the size of mapped regions is a multiple of system page size
 - mapping regions which are way smaller than that can result in a
- significant waste of memory
- memory mapping must fit in the process address space
 - on 32 bits systems, a large number of mappings of various sizes might result in memory fragmentation
 - it then becomes harder to find continuous space to grant large memory mappings
 - the problem is substantially diminished on 64 bits systems
- there is kernel overhead in maintaining mappings
 - for small mappings, the overhead can dominate the advantages
 - memory mapped I/O is best used with large files and random access

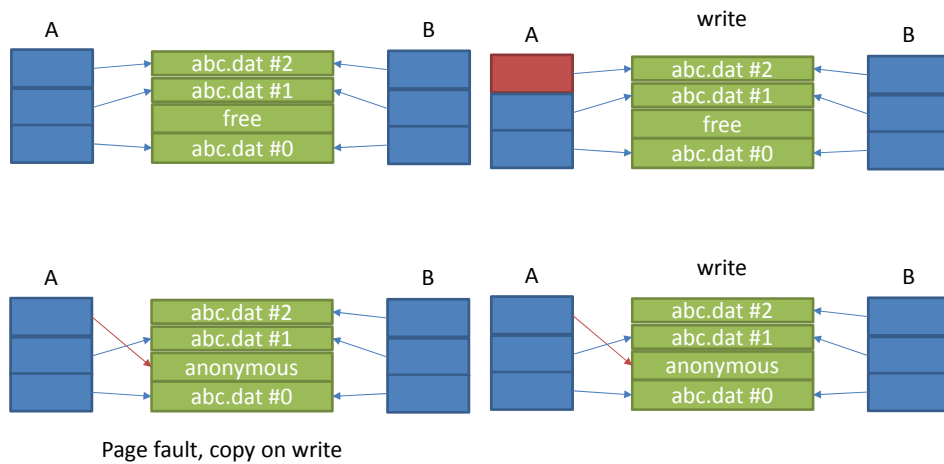
Ref: Programmation Systèmes /Memory Mapping, Stefano Zacchiroli

Page Cache

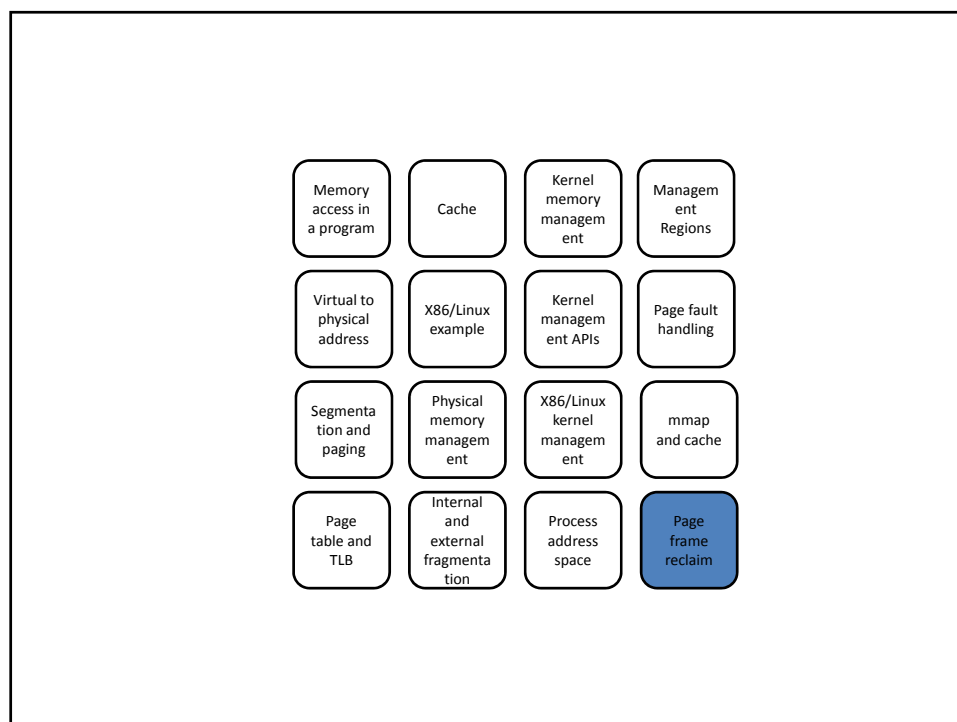
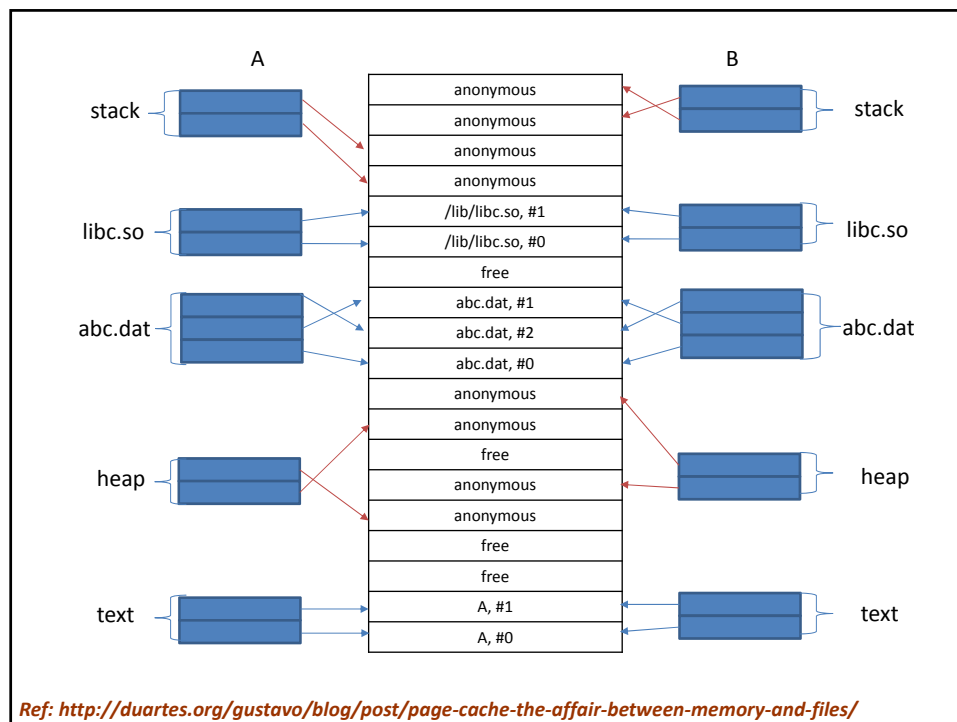


Ref: <http://duartes.org/gustavo/blog/post/page-cache-the-affair-between-memory-and-files/>

Read-only
Read/write



Ref: <http://duartes.org/gustavo/blog/post/page-cache-the-affair-between-memory-and-files/>



The Types of Pages Considered by the PFRA

- *Ref: Understanding the Linux Kernel, Daniel P. Bovet, Marco Cesati*

- *Ref: CS161: Operating Systems, Matt Welsh*

Implementation of PFRA

- *Ref: Understanding the Linux Kernel, Daniel P. Bovet, Marco Cesati*